Advanced Graph Theory Notes

Sam Spiro

January 20, 2026

Contents

0 Overview and Basic Definitions

These notes are intended for a graduate course in graph theory which assumes the reader is already familiar with basic graph theory terms and definitions (see also Section 0.1 for a recap of these definitions). **You should expect many typos and missing references**.

The first half of these notes centers on two of the main areas of modern graph theory: extremal graph theory and structural graph theory. Broadly speaking, extremal graph theory ask questions of the form: how "large" can a graph be if it satisfies a certain property? Structural graph theory, on the other hand, broadly speaking aims to characterize families of graphs which satisfy a certain property. It is worth noting that the exact line between these two areas is rather vague, so some topics may have crossover between each other. It should also be said that I am an extremal graph theorist, so there will certainly be a bias these topics.

The second half of the text centers around "bonus" material which delves into specific methods for solving graph theory problems, as well as auxiliary topics which could be entire courses on their own.

0.1 Very Basic Graph Theory Definitions

Here we briefly recall the basic definitions and notations for graphs that we use throughout the text.

The Essentials:

- A *graph* G is a pair of sets (V, E) with E a set of 2-element subsets of V, i.e. $E \subseteq \{\{x, y\} : x, y \in V,\ x \neq y\}$. The set V is called the *vertex set* of G and its elements are called *vertices*, while the set E is called the *edge set* of G and its elements are called *edges*. We will typically denote edges $\{x, y\}$ by the simpler notation xy.

  Eg $(\{1, 2, 3, 4\}, \{12, 23, 13, 14\})$ is a graph. Often it's easier to depict graphs by pictures (and how exactly we draw the picture doesn't matter).

- Throughout this text we will only consider finite graphs, ie graphs with $|V| < \infty$, though we emphasize that interesting things can be said regarding infinite graphs.

- Throughout this text we will almost always work with graphs without repeated edges (ie E is a set rather than a multiset) and graphs without oriented edges (ie each edge is an *unordered* pair of vertices, meaning $xy = yx$).

- We will often write $V(G)$ and $E(G)$ to denote the vertex and edge sets of a graph G, and we write $v(G) = |V(G)|$ and $e(G) = |E(G)|$.

- We say two vertices x, y are *adjacent* or *neighbors* if $xy \in E(G)$, and we sometimes denote this be writing $x \sim y$.

- Given a vertex x we define the *neighborhood* of x by $N(x) = \{$vertices that are adjacent to x in $G\}$. We define the *degree* of x by $\deg(x) = |N(x)|$. Whenever the graph G is not clear from context we will write $N_G(x)$ and $\deg_G(x)$.

- We say that a graph $G' = (V', E')$ is a *subgraph* of another graph $G = (V, E)$ if $V' \subseteq V$ and $E' \subseteq E$. In this case we write $G' \subseteq G$.

- We say two graphs G, H are isomorphic if there exists a bijection $\phi : V(G) \to V(H)$ such that $x, y \in V(G)$ are adjacent in G if and only if $\phi(x), \phi(y) \in V(H)$ are adjacent in H for all x, y.

Paths and Connectivity:

- A *path* in a graph G is sequence of distinct adjacent vertices $(x_1, x_2, \ldots, x_t)$, and we say such a path is a path from x_1 to x_t and that it has *length* $t-1$ (i.e. the length of the path is the number of edges it has).

- A graph is *connected* if for any two pair of vertices there exists a path from x to y.

- The *distance* between two vertices x, y, denoted $\mathrm{dist}(x, y)$, is the length of the shortest path from x to y (with $\mathrm{dist}(x, y) = \infty$ if no such path exists).

Graph Operations and Subgraphs

- Given a set S and an integer k, we let $\binom{S}{k}$ denote the set of all subsets of S of size k. For example, our definition of a graph is equivalent to saying that $E \subseteq \binom{V}{2}$.

- Given a graph G we define its *complement* $\overline{G}$ to be the graph obtained by replacing all edges with non-edges and vice versa. That is, $\overline{G}$ is the graph with vertex set $V(G)$ and edge set $\binom{V(G)}{2} \setminus E(G)$.

- Given a graph G and a set of vertices $S \subseteq V(G)$, we define $G - S$ to be the graph obtained by deleting S and all edges incident to it. That is, $V(G - S) = V(G) \setminus S$ and $E(G-S) = E(G) \setminus \{e : e \cap S \neq \emptyset\}$. If $S = \{x\}$ then we will denote this simply by $G - x$. Similarly if xy is an edge of G we define $G - xy$ to be the graph obtained by deleting the edge xy.

- A subgraph $G' \subseteq G$ is said to be *induced* if it is of the form $G - S$ for some set of vertices S. Given a set of vertices V we will sometimes write $G[V]$ to be the induced subgraph with vertex set V, i.e. $G[V] = G - V(G) \setminus V$.

- A subgraph $G' \subseteq G$ is called *spanning* if $V(G') = V(G)$.

Independent Sets and Colorings

- A set of vertices I is *independent* if no two vertices $x, y \in I$ are adjacent to each other.
- A graph is bipartite if there exists a partition of $V(G)$ into two independent sets.
- Given a graph G and an integer k, a *proper k-coloring* is a map $\phi : V(G) \to [k]$ with the property that adjacent vertices $x, y \in V(G)$ have $\phi(x) \neq \phi(y)$. The smallest k for which G has a proper k-coloring is called the *chromatic number* of G and is denoted $\chi(G)$.

Forests and Trees

- A graph is a *forest* if it contains no cycles (i.e. no subgraph isomorphic to a cycle graph C_ℓ). A *tree* is a forest which is connected.
- A vertex of degree 0 is called an *isolated vertex*. A vertex of degree 1 (especially in the context of trees and forests) is called a *leaf*.

0.2 Common Graph Families and Parameters

We record notation for graphs that will appear throughout the text.

- K_n denotes the n-vertex complete graph, i.e. the unique n-vertex graph with all $\binom{n}{2}$ edges.
- $K_{s,t}$ denotes the complete bipartite graph which has s vertices in one part and t vertices in the other.
- C_ℓ denotes the cycle graph of length ℓ.
- P_r denotes the path graph with r vertices (NOTE: some authors would denote this by P_{r-1}).

We record notation for graph parameters that will appear throughout the text, where here G denotes an arbitrary graph.

- $\delta(G)$ is the minimum degree of G, i.e. $\delta(G) = \min_{x \in V(G)} \deg(x)$.
- $\Delta(G)$ is the maximum degree of G, i.e. $\Delta(G) = \max_{x \in V(G)} \deg(x)$.
- $\alpha(G)$ is the independence number of G, which is the largest size of an independent set of G.
- $\chi(G)$ is the chromatic number of G, which is the smallest integer k such that G has a proper k-coloring.

0.3 Asymptotic Notation

Eventually in the text it will be convenient for us to make use of the following asymptotic notation which we record here for ease of reference. We emphasize that this notation will be redefined when it first appears in the text, so there is no need to memorize this right now.

Let $f(n), g(n)$ be two functions.

- We write $f(n) = O(g(n))$ if there exists a constant $C > 0$ such that $f(n) \leq Cg(n)$ for all n.
- We write $f(n) = \Omega(g(n))$ if there exists a constant $c > 0$ such that $f(n) \leq cg(n)$ for all n.
- We write $f(n) = \Theta(g(n))$ if $f(n) = O(g(n))$ and $f(n) = \Omega(g(n))$. In this case we say that f, g have the same *order of magnitude*.
- We write $f(n) \sim g(n)$ if $\lim_{n\to\infty} \frac{f(n)}{g(n)} = 1$. In this case we say that f, g are *asymptotic* to each other.
- We write $f(n) = o(g(n))$ if $\lim_{n\to\infty} \frac{f(n)}{g(n)} = 0$. In particular, writing $f(n) = o(1)$ means $\lim_{n\to\infty} f(n) = 0$.

0.4 Inequalities

Many proofs in extremal combinatorics rely on basic inequalities from analysis. Here we record the most important of these that we will use.

Theorem (Cauchy-Schwarz Inequality). *If $x_1, \ldots, x_n$ and $y_1, \ldots, y_n$ are real numbers, then*

$$\sum_{i=1}^{n} x_i y_i \leq \left(\sum_{i=1}^{n} x_i^2\right)^{1/2} \left(\sum_{i=1}^{n} y_i^2\right)^{1/2}$$

Personally, we like to remember the statement of Cauchy-Schwarz by noting that it follows from the vector equality $\langle \mathbf{x}, \mathbf{y} \rangle = \cos\theta \, \|\mathbf{x}\| \, \|\mathbf{y}\|$ where θ is the angle between the vectors $\mathbf{x}, \mathbf{y}$.

For the next inequality, recall that a function $\phi : \mathbb{R} \to \mathbb{R}$ is *convex* if for all $0 \leq t \leq 1$ and $x, y \in \mathbb{R}$ we have $\phi(tx + (1-t)y) \leq t\phi(x) + (1-t)\phi(y)$.

Theorem (Jensen's Inequality). *If ϕ is a convex function and $x_1, \ldots, x_n \in \mathbb{R}$, then*

$$\sum_{i=1}^{n} \phi(x_i) \geq n\phi\left(n^{-1} \sum_{i=1}^{n} x_i\right).$$

That is, this sum is minimized when each x_i is equal to their average $n^{-1} \sum x_{i'}$. We note for later that for any integer $t \geq 1$, the function $\binom{x}{t} := \frac{x(x-1)\cdots(x-t+1)}{t!}$ is convex.

0.5 Exercises

Each chapter will end with a set of exercises. Following the notation of Stanley, we will add numbers after each exercise to indicate the problem's rough level of difficulty as follows:

- [1] problems are elementary and routine requiring little to no thought,
- [2] problems have simple solutions (though that does not necessarily mean it is easy to find such a solution!),
- [3] problems tend to have involved solutions,
- [4] problems have extremely difficult solutions (to the extent that such questions should never be used in a classroom setting),
- [5] problems are unsolved open problems.

Additionally, plus and minus symbols may be used to indicate higher or lower levels of difficulty for the problem. For example, a [2+] problem might have a simple solution that's pretty challenging to find, while a [3-] problem might have an involved solution that's actually not too hard to work out. Ultimately, all of the ratings that I give are only rough estimates and the reader may find a given [3] problem easier to solve than a [2-] depending on the circumstances.

With that preamble out of the way, we begin with some "elementary" (though not necessarily easy) graph theory problems.

1. (Handshaking Lemma) Prove that every graph G has $\sum_{x\in V(G)} \deg(x) = 2e(G)$ [2-].

2. Prove that every graph G with $v(G) \geq 2$ contains two vertices with the same degree [2-].

3. Prove that for every graph G, either G or its complement $\overline{G}$ is connected [2-].

4. Prove that a graph is bipartite if and only if it contains no odd cycles [2-].

5. Prove that for every graph G, the set of edges $E(G)$ can be partitioned into cycles if and only if every very vertex of G has even degree [2+].

* * *

6. Recall that a graph is d-regular if $\deg(u) = d$ for every vertex u. Prove for all integers $0 \leq d < n$ that there exists an n-vertex d-regular graph if and only if at least one of d or n is even [2].

7. A graph is said to have girth g if it contains a cycle of length g and no cycles of shorter length.

(a) Prove that for all integers $d, g \geq 2$, there exists a d-regular graph of girth g [2+].

(b) Prove that if G is a d-regular graph of girth g, then

$$v(G) \leq ???.$$

[2]

(c) Show that the bound above is tight for $d = 3, g = 5$ [1+].

* * *

8. Prove that $\chi(G)\alpha(G) \leq v(G)$ for all graphs G [2-].

9. Prove that $\alpha(G) \geq \frac{v(G)}{\Delta(G)+1}$ for all graphs G [2].

10. Prove that if a graph G is triangle-free (i.e. if G contains no subgraph isomorphic to K_3) then $\alpha(G) \geq \sqrt{v(G)}$ [2-].

* * *

11. Prove that every tree T with $v(T) \geq 2$ has at least two leaves.

12. Prove that for every tree T, there exists an ordering of its vertices $v_1, \dots, v_n$ such that for all $2 \leq i \leq n$, there exists an integer j_i such that $N(v_i) \cap \{v_1, \dots, v_{i-1}\} = \{v_{j_i}\}$ [1+].

13. Prove various characterizations of trees

14. (Helly Theorem for Trees) Let T be a tree and $\mathcal{T}$ a set of subtrees of T (i.e. a set of subgraphs of T which are themselves trees). Prove that if $V(T') \cap V(T'') \neq \emptyset$ for all $T', T'' \in \mathcal{T}$, then there exists a vertex $v \in \bigcap_{T' \in \mathcal{T}} V(T')$ [2+].

Part I

Extremal Graph Theory

As mentioned in the introduction, extremal graph theory broadly speaking asks questions of the form: how "large" can a graph be if it satisfies a certain property?

What exactly "large" means depends on the type of problem one is considering, with some popular choices being the number of edges, the number of vertices, and the minimum degree of the graph in question. Each of these choices (together with an appropriate choice of "property") gives rise to three of the main topics of extremal graph theory: Turán problems, Ramsey problems, and Dirac problems; see the table below for a brief outline. Each of these types of problems will be the main topic of focus for the forthcoming chapters.

Measurement		Property		Type of Problem
Number of edges	$+$	Triangle-free	$=$	Turán Problems: Section 1
Number of vertices	$+$	G and $\overline{G}$ are triangle-free	$=$	Ramsey Problems: Section 3
Minimum degree	$+$	non-Hamiltonian	$=$	Dirac Problems: Section 2

Figure 1: A table of measures of "largeness", properties that one can consider, and the problems that these produce. Note that in each case, the given property is harder to fulfill the "larger" G is with respect to its measurement, which is a hallmark of a good extremal problem.

1 Forbidden Subgraphs and Turán Problems

Turán Problems broadly ask: how many edges can an n-vertex graph have if it does not contain a copy of a given graph F? Specifically, the we will work with the following throughout this chapter.

Definition 1. Given two graphs F, G, we say that G is *F-free* if G does not contain a subgraph which is isomorphic to F. Given an integer $n \geq 1$, we define the *Turán number* or *extremal number* $\mathrm{ex}(n, F)$ to be the maximum number of edges that an n-vertex F-free graph can have.

The name of the game now is to try and either determine or bound $\mathrm{ex}(n, F)$ for various choices of F.

1.1 Forbidding C_4 and Complete Bipartite Graphs

Perhaps the first question we need to answer is: why should we care about Turán problems in the first place? There are many possible answers to this question, here are a few of my own personal reasons:

- They are natural extremal problem to consider.
- They have applications to various areas of mathematics.
- Solutions to Turán problems often use cool and deep results from other areas of mathematics in interesting ways.
- They're fun!

To try and illustrate these points above, we will begin by studying the Turán problem for $F = C_4$. Historically, this is the second Turán problem to be considered (we will look at the first problem in the following section) and was largely solved by Erdős in year due to its connection to a certain problem in number theory.

The Upper Bound. We begin by establishing an upper bound for this Turán number.

Theorem 1.1. *We have*

$$\mathrm{ex}(n, C_4) \leq \frac{n\sqrt{4n-3}+n}{4}.$$

That is, every n-vertex C_4-free graph has at most this many edges.

We emphasize that this is not a very pretty looking upper bound; we will address this further shortly after the proof.

Proof. In order to prove any upper bound for this problem, we need to get some understanding of what it means for a graph to be C_4-free graph. After thinking about it for long enough, one might come up with the following observation: a graph is C_4-free if and only if every pair of distinct vertices u, v has at most one common neighbor, i.e. there is at most one vertex in $N(u) \cap N(v)$. Indeed, the existence of two vertices in this set together with u, v would exactly define a C_4 in our graph.

Now, a priori, it is not immediate how to use the fact that pairs of vertices have at most one common neighbor to bound the number of edges in our graph. However, one can use it to bound the number of some other object which is "almost" an edge. Namely, let

$$\mathcal{P} = \{(\{u, v\}, x) \in V(G)^3 : u \sim x \sim v,\ u \neq v\},$$

which essentially just encodes the set of P_3's in G. Note that each element of $\mathcal{P}$ can be uniquely identified by picking two distinct vertices to play the roles of u, v together with a common neighbor of these vertices to play the role of x. As such, we have

$$|\mathcal{P}| = \sum_{u \neq v} |N(u) \cap N(v)| \leq \sum_{u,v} 1 = \binom{n}{2},$$

with the inequality using that our graph is C_4-free. Now, we got the first equality above by identifying each element of $\mathcal{P}$ by its first and last vertices u, v and then picking some common neighbor x. Alternatively, we could identify each element of $\mathcal{P}$ by specifying its middle vertex x together with two distinct neighbors u, v of x. As such, we also have

$$|\mathcal{P}| = \sum_{x \in V(G)} \binom{\deg(x)}{2} \geq n \binom{n^{-1} \sum_x \deg(x)}{2} = n \binom{n^{-1} \cdot 2e(G)}{2},$$

where this inequality used Jensen's inequality together with the fact that $\binom{a}{2}$ is a convex function, and the last equality used that $\sum_x \deg(x) = 2e(G)$. Comparing this to the upper bound for $|\mathcal{P}|$ we found above gives

$$n \binom{n^{-1} \cdot 2e(G)}{2} \leq \binom{n}{2}, \tag{1}$$

or equivalently

$$(2e(G))(2n^{-1}e(G) - 1) \leq n(n-1).$$

This in turn is equivalent to having

$$4e(G)^2 - 2ne(G) - n^2(n-1) \leq 0,$$

and solving this exactly gives the desired bound on $e(G)$.

Somewhere in the text I should call this a double counting argument and maybe mention the word cherries/P_2.

□

While the bound of Theorem 1.1 is truly the best we can do using our approach, it is often not a good idea in extremal combinatorics to do things so precisely.

Mantra 1. It is often better to use (slightly) "wastefull" bounds in extremal combinatorics to have cleaner proofs and theorem statements.

Knowing when exactly and how to derive such "crude" bounds is an important skill to have in extremal combinatorics, since in practice we do not know a priori if the approach we are currently playing around with is going to give something useful in the end, and until that point it is a bad idea to harp over minute details in the argument.

For example, let us consider the point in the proof where we reached (1). Here an expert might simplify their lives by observing that simple inequalities for binomial coefficients yield

$$n \cdot \frac{1}{2}(n^{-1}2e(G) - 1)^2 \leq n\binom{n^{-1} \cdot 2e(G)}{2} \leq \binom{n}{2} \leq \frac{1}{2}n^2,$$

and rearranging this gives

$$n^{-1}2e(G) - 1 \leq n^{1/2},$$

and hence

$$e(G) \leq \frac{1}{2}n^{3/2} + n.$$

Note that this is extremely close to the optimal bound we get in Theorem 1.1. In particular, one can show that both bounds are ultimately of the form $\mathrm{ex}(n, C_4) \leq \frac{1}{2}n^{3/2} + Cn$ for some sufficiently large constant C. This means that our weakening above captures the "main part" of the bound from Theorem 1.1, in the sense that for n very large the two numbers are very close to each other.

It will be useful going forward to develop notation to measure more precisely what exactly we mean by "very close to each other".

Definition 2. Let $f(n), g(n)$ be two functions.

- We write $f(n) = O(g(n))$ if there exists a constant $C > 0$ such that $f(n) \leq Cg(n)$ for all n. In particular, our remark in the paragraph above is equivalent to saying that our two bounds give[1] $\mathrm{ex}(n, C_4) \leq \frac{1}{2}n^{3/2} + O(n)$.

- We write $f(n) = \Omega(g(n))$ if there exists a constant $c > 0$ such that $f(n) \leq cg(n)$ for all n. Whenever we write this, we will often implicitly assume that we consider n large enough so that $f(n) > 0$. For example, if we write $\mathrm{ex}(n, F) = \Omega(1)$ we will implicitly be assuming $n \geq 2$.s

- We write $f(n) = \Theta(g(n))$ if $f(n) = O(g(n))$ and $f(n) = \Omega(g(n))$. In this case we say that f, g have the same *order of magnitude*.

- We write $f(n) \sim g(n)$ if $\lim_{n\to\infty} \frac{f(n)}{g(n)} = 1$. In this case we say that f, g are *asymptotic* to each other.

[1] A very persnickety reader might object that actually this doesn't exactly agree with the definition given: the real thing that should be written is $\mathrm{ex}(n, C_4) - \frac{1}{2}n^{3/2} = O(n)$ and the "algebra" of moving $\frac{1}{2}n^{3/2}$ to the other side is not actually valid. It is, however, common practice in the field to use these somewhat imprecise notational implementations in order to make statements easier to read and write, which is the ultimate goal of introducing this in the first place.

- We write $f(n) = o(g(n))$ if $\lim_{n\to\infty} \frac{f(n)}{g(n)} = 0$. In particular, writing $f(n) = o(1)$ means $\lim_{n\to\infty} f(n) = 0$.

Applications. Theorem 1.1 has a number of applications to other areas of mathematics. We will consider one quick example from discrete geometry.

Let $\mathcal{P}$ be a set of points of $\mathbb{R}^2$ and let $\mathcal{L}$ be a set of lines in $\mathbb{R}^2$. We say that a point $p \in \mathcal{P}$ and a line $\ell \in \mathcal{L}$ are *incident* if p lies on the line ℓ. We let $I(\mathcal{P}, \mathcal{L})$ to denote the number of pairs $(p, \ell) \in \mathcal{P} \times \mathcal{L}$ with p and ℓ incident. A natural extremal question to ask is: what is the maximum number of incidences that a given number of points and line can obtain? Trivially one can do no better than n^2, but it is not so immediate how to improve this. We will be able to achieve such an improvement using our Turán result Theorem 1.1.

Corollary 1.2. *If $\mathcal{P}$ is a set of n points in $\mathbb{R}^2$ and if $\mathcal{L}$ is a set of n lines in $\mathbb{R}^2$, then*

$$I(\mathcal{P}, \mathcal{L}) = O(n^{3/2}).$$

Proof. As is often the case for applications, we begin by defining an auxiliary graph related to our problem at hand. To this end, define a bipartite graph G whose vertex set is $\mathcal{P} \cup \mathcal{L}$ where we have $p \sim \ell$ if and only if p and ℓ are incident. Observe that $I(\mathcal{P}, \mathcal{L}) = e(G)$, so bounding the number of incidences is exactly the same thing as bounding the number of edges of G.

Now, for arbitrary bipartite graphs G we could of course have $e(G)$ as large as n^2, but we have some additional structure to work with because G is coming from a set of points and lines. In particular, because every pair of lines intersect in at most one point, G can not contain a C_4 (since such a subgraph would consist of vertices p_1, p_2, ℓ_1, ℓ_2 with p_1, p_2 points common to both ℓ_1 and ℓ_2). This together with the fact that $v(G) = |\mathcal{P}| + |\mathcal{L}| = 2n$ immediately implies that

$$I(\mathcal{P}, \mathcal{L}) = e(G) \le \mathrm{ex}(2n, C_4) = O((2n)^{3/2}) = O(n^{3/2}),$$

with this last step using that this "big oh" notation is not affected by multiplying by a fixed constant. $\square$

While it is neat that we could obtain this purely geometric result using graph theory, we should note that the bound of Corollary 1.2 is not tight, and in fact the true bound is $I(\mathcal{P}, \mathcal{L}) = O(n^{4/3})$. The fact that we obtained a subpar bound should perhaps not come as a surprise, as we used almost no information about the geometry of the Euclidean plane $\mathbb{R}^2$ in our argument. It is, however, possible to derive this optimal bound of $O(n^{4/3})$ if one uses Theorem 1.1 together with some appropriate geometric tools (such as real polynomial partiontionings). We will not go into this here, but see eg the book by Sheffer for a lot more on this problem and more.

The Lower Bound. Theorem 1.1 shows that $\mathrm{ex}(n, C_4) = O(n^{3/2})$. The immediate question is: is this tight? This is an important question for us to figure out, since e.g. any improvement to Theorem 1.1 would give an improvement to our bound in Corollary 1.2 as well as to any other application we can come up with for $\mathrm{ex}(n, C_4)$.

To see whether our bound is tight, we need to prove a lower bound for $\mathrm{ex}(n, C_4)$, i.e. to construct n-vertex graphs with many edges and no C_4's. This, as the reader is welcome to try for themselves, is not so easy to do. To make some headway on this, we usethe following mantra.

Mantra 2. To find a lower bound construction for extremal problems, we should ask ourselves what would need to happen for our extremal upper bound to be (exactly) sharp.

In our case we ask: what would need to happen for us to have $\mathrm{ex}(n, C_4) = \frac{n\sqrt{4n-3}+n}{4}$? Well, this would happen precisely if every inequality throughout our proof of Theorem 1.1 were in fact an *equality*. In particular, our very first inequality $\sum_{u \neq v} |N(u) \cap N(v)| \leq \binom{n}{2}$ must be an equality, and this would imply that *every* pair of distinct vertices in G has exactly 1 common neighbor. Now we have to ask...is this ever possible?

Well, if you think about it for long enough, you might have the wild idea that "every two vertices has exactly 1 common neighbor" is kind of analogous to the statement "every two points in $\mathbb{R}^2$ lie on exactly one common line." Riffing off of this as well as what we did for our application in Corollary 1.2, what if we defined a bipartite graph G by taking a set of points $\mathcal{P}$ and a set of lines $\mathcal{L}$ and making a point p adjacent to a line ℓ if and only if they are incident? Such a graph will automatically be C_4-free due to the geometry of the situation, so we will win if we can find some points and lines with many incidences.

As hinted at just after Corollary 1.2, it is possible to find n points and lines in $\mathbb{R}^2$ such that $I(\mathcal{P}, \mathcal{L}) = \Omega(n^{4/3})$, giving a corresponding lower bound to $\mathrm{ex}(n, C_4)$, but this is as good as we can hope to do in Euclidean space. However, another wild thought based on what we said around Corollary 1.2 is that our idea of using points and lines does not fundamentally rely on the full geometry of Euclidean space: we only needed the very basic property that two points line on at most one line, and such a property holds for many different types of geometries. In particular, since we're working with finite graphs...why not try and do something with geometries over finite fields?

Recall from algebra[2] that for every prime power q there exists a field $\mathbb{F}_q$ of order q. Again going off what we did in Euclidean space, we want to consider a set of points and lines from the plane $\mathbb{F}_q^2 = \{(x, y) : x, y \in \mathbb{F}_q\}$. There might be some particularly clever choices of points and lines that we could make here, but since we are just playing around, why don't we go ahead and just take all of them. That is, we will take $\mathcal{P} = \mathbb{F}_q^2$ and $\mathcal{L}$ all of the lines in $\mathbb{F}_q^2$. To be clear, lines in $\mathbb{F}_q^2$ are just sets of points in $\mathbb{F}_q^2$ taking on one of two forms: for $a, b \in \mathbb{F}_q$ we define the line $\ell_{a,b} = \{(x, ax + b) : x \in \mathbb{F}_q\}$, and for $c \in \mathbb{F}_q$ we define the vertical lines $\ell_c = \{(c, y) : y \in \mathbb{F}_q\}$. Now define a bipartite graph G_q on $\mathcal{P} \cup \mathcal{L}$ where $p \sim \ell$ if and only if $p \in \ell$. We leave it as an exercise to the reader to verify that G_q is indeed C_4-free. To count $e(G_q)$, we observe that the total number of lines is $q^2 + q$ and that each line is incident to exactly q points, and as such $e(G_q) = q^3 + q^2$. Because the total number of vertices in G_q is exactly $2q^2 + q$, we in total conclude for any prime power q that

$$\mathrm{ex}(2q^2 + q, C_4) \geq q^3 + q^2.$$

By considering $n = 2q^2 + q \approx 2q^2$ or equivalent $q \approx (n/2)^{1/2}$, we find that for infinitely many integers n that $\mathrm{ex}(n, C_4)$ is at least $q^3 \approx (n/2)^{3/2} = 2^{-3/2} n^{3/2}$. As such, the upper bound of $\mathrm{ex}(n, C_4) = O(n^{3/2})$ really is the best we can do for general n! In fact, some basic number theory facts let us prove the following.

Theorem 1.3. *We have* $\mathrm{ex}(n, C_4) = \Theta(n^{3/2})$.

[2] Any reader scared of algebra should be reassured that this is the only fact you need to recall from algebra.

Proof. By Theorem 1.1 we have for n large enough that, say, $\mathrm{ex}(n, C_4) \le n^{3/2}$, proving $\mathrm{ex}(n, C_4) = O(n^{3/2})$.

Now consider any integer $n \ge 12$. By Bertrend's postulate, there exists a prime number p with $\frac{1}{2}\sqrt{n/3} \le p \le \sqrt{n/3}$. This in particular implies $n \ge 3p^2 \ge 2p^2 + p$, which together with our discussion above implies

$$\mathrm{ex}(n, C_4) \ge \mathrm{ex}(2p^2 + p, C_4) \ge p^3 \ge (12)^{-3/2} n^{3/2},$$

proving that $\mathrm{ex}(n, C_4) = \Omega(n^{3/2})$ and hence that $\mathrm{ex}(n, C_4) = \Theta(n^{3/2})$ as desired. □

We personally find it fascinating that one can use ideas from algebra and geometry to solve the purely combinatorial problem of determining $\mathrm{ex}(n, C_4)$. This is in fact a very common phenomenon.

Mantra 3. To solve a combinatorics problem, one often needs ideas and tools from other areas of math. As such, any extra knowledge you have outside of combinatorics is always useful to keep in the back of your mind!

This mantra is intended to be inspirational rather than intimidating. In particular, even if you don't have hardly any knowledge in areas outside of combinatorics (such as myself), you can still make it very far, its just that some problems in particular may elude your grasps until you figure out the right tool needed to crack it.

Even Better Lower Bounds. We've done pretty good so far with our lower bounds for $\mathrm{ex}(n, C_4)$, but we can go even farther.

Mantra 4. Once you prove something, see if you can prove something even better.

In particular, given that we have determined the order of magnitude $\mathrm{ex}(n, C_4) = \Theta(n^{3/2})$, we should next ask ourselves if we can prove that $\mathrm{ex}(n, C_4) \sim cn^{3/2}$ for some constant c. We emphasize that doing this will require a bit more algebra/geometry than before, and as such the reader may wish to skip over this part of the text if they're already overwhelmed.

Returning back to the problem at hand, we know up to this point (at least for certain values of n) that

$$2^{-3/2} n^{3/2} + o(n^{3/2}) \le \mathrm{ex}(n, C_4) \le 2^{-1} n^{3/2} + o(n^{3/2}),$$

and we need to figure out if we can sharpen either of these bounds. For this, it is useful to analyze "why" our lower bound proof does not match the bound we got in the upper bound. After all, in our construction every pair of points really does have exactly one common neighbor. However, if we look back at what motivated our construction in the first place, we recall that for the upper bound for Theorem 1.1 to be exactly sharp that we need every pair of *vertices* to have a common neighbor, and there is no hope of that happening for our current graph because G_q is bipartite (meaning a given point and a given line will never have any common neighbors in G_q).

It is not so immediate how to fix this problem, as the underlying motivation for our construction relied on working with both points and lines which intrinsically are different objects from each other. But, if we stare at things long enough, we might realize that our lines $\ell_{a,b}$ are indexed by points in $\mathbb{F}_q^2$, and as such, one might possibly have the idea where we could consider a graph G where its vertex set is just $\mathbb{F}_q^2$ but where a point (x,y) corresponds to both the point itself and the line $\ell_{x,y}$. That is, we want to define a graph on $\mathbb{F}_q^2$ where $(x,y) \sim (a,b)$ if and only if $(x,y) \sim \ell_{a,b}$. While this is a noble idea, an immediate issue in this definition is that this edge relation is not symmetric. That is, having $(x,y) \in \ell_{a,b}$ does not imply $(a,b) \in \ell_{x,y}$ (i.e. $y = ax + b$ does not mean $b = xa + y$). At a very high level the issue here with the idea of identifying points with a corresponding line is that points and lines are not truly "dual" to each other in $\mathbb{F}_q^2$. However, this can be fixed by going to yet another type of geometry, namely projective geometry.

Insert better intuition on projective geometries at some point.

To define things, consider the set of triples $T = \{(x,y,z) : x,y,z \in \mathbb{F}_q^3\} \setminus \{(0,0,0)\}$ and define an equivalence relation (not to be confused with an edge relation) by having $(x,y,z) \equiv (\alpha x, \alpha y, \alpha z)$ for all $\alpha \in \mathbb{F}_q \setminus \{0\}$. Let $[x,y,z]$ denote the equivalence class containing (x,y,z), and define our set of "points" $\mathcal{P}$ to be the set of all such equivalence classes. For each $[a,b,c] \in \mathcal{P}$ we define the line $\ell_{[a,b,c]} = \{[x,y,z] : ax + by + cz = 0\}$. Note that this definition is well-defined (i.e. it does not matter whether we write $[x,y,z]$ or $[\alpha x, \alpha y, \alpha z]$) since having $ax + by + cz = 0$ implies $\alpha ax + \alpha by + \alpha cz = 0$ for all $\alpha \neq 0$. Also note that this definition is truly "dual" in points and lines, in that $[x,y,z] \in \ell_{[a,b,c]}$ if and only if $[a,b,c] \in \ell_{[x,y,z]}$. Motivated by this and our ideas from above, we define a graph G_q^* on $\mathcal{P}$ where $[x,y,z] \sim [a,b,c]$ if and only if $[x,y,z] \in \ell_{[a,b,c]}$. We leave it as an exercise to verify that G_q^* is C_4-free, that $v(G_q^*) = q^2 + q + 1$, and that $e(G_q^*) = \frac{1}{2}(q+1)(q^2+q+1)$.

Similar to before, if we take $n = q^2 + q + 1 \approx q^2$, then we see that this shows $\mathrm{ex}(n, C_4)$ is at least $\frac{1}{2}q^3 \approx \frac{1}{2}n^{3/2}$, exactly matching the asymptotic bound from Theorem 1.1! Actually, even more is true: one can check that the upper bound $\frac{n\sqrt{4n-3}+n}{4}$ is actually *exactly* tight in this case. That is, for all prime powers q, we have

$$\mathrm{ex}(q^2 + q + 1, C_4) = \frac{1}{2}(q+1)(q^2+q+1).$$

Generalizations. Given our success with studying the Turán problem for C_4, we should go on and ask to what extent can the ideas here be used to prove bounds for other graphs F. Naively one might first consider the problem for other cycles C_ℓ, but this turns out to be pretty difficult. Instead, the "correct" generalization of the ideas we have here are for complete bipartite graphs $K_{s,t}$ in general beyond just that of $K_{2,2} = C_4$. For example, we leave it as an exercise to generalize the upper bound in Theorem 1.1 to prove the following general upper bound.

Theorem 1.4 (Kővári-Sós-Turán Theorem)**.** *For all integers $s,t \geq 1$, we have*

$$\mathrm{ex}(n, K_{s,t}) = O_{s,t}(n^{2-1/s}).$$

Here we add the s,t subscript to the big-oh notation to emphasize that the implicit constant depends on s,t. This is not entirely necessary since we fix s,t at the start of the theorem, but it is sometimes nice to emphasize this for clarity.

This gives an upper bound, what about a corresponding lower bound? Our lower bound $\mathrm{ex}(n, C_4) = \Omega(n^{3/2})$ immediately implies $\mathrm{ex}(n, K_{2,t}) = \Omega(n^{3/2})$ for all $t \geq 2$, giving the correct order of magnitude. In fact, Füredi**REF**improved the lower bound for $\mathrm{ex}(n, K_{2,t})$ even further, giving a tight asymptotic bound. With some effort, one can generalize the geometric intuition we had for C_4 to prove $\mathrm{ex}(n, K_{3,t}) = \Theta(n^{5/3})$ for all $t \geq 3$, roughly be replacing the intuition of "two lines intersect in at most one point" with "three spheres intersect in at most two points." Despite this success, the next case of this problem remains open.

Open Problem 1.5. *Determine the order of magnitude of* $\mathrm{ex}(n, K_{4,4})$.

Similarly $\mathrm{ex}(n, K_{s,s})$ remains open for all $s \geq 4$. However, it turns out that we can solve this problem for $K_{s,t}$ whenever t is sufficiently large in terms of s.

Theorem 1.6. *For all $s \geq 2$, there exists an integer t_0 such that $\mathrm{ex}(n, K_{s,t}) = \Theta(n^{2-1/s})$ for all $t \geq t_0$.*

The first result of this form was proven by Authors who showed one can take $t_0 =$Something by using an explicit algebraic construction like we had for G_q^*. The best current bound is due to Bukh who recently showed one can take $t_0 = 9^{(1+o(1))s}$ by using a *random* algebraic construction.

1.2 Forbidding Cliques

Now that we've all been convinced that studying $\mathrm{ex}(n, F)$ is an interesting problem, we need to figure out some graphs F for which we can effectively bound (or even determine) $\mathrm{ex}(n, F)$. As a starting step, we can think about this problem for small graphs F. A moment's thought shows that it is quite easy to determine $\mathrm{ex}(n, F)$ for every graph F with $v(F) \leq 3$ *except* for the graph $F = K_3$, which is the smallest non-trivial instance of this problem. The full solution to this problem is a classical result of Mantel from 1907.

Theorem 1.7 (Mantel's Theorem)**.** *We have $\mathrm{ex}(n, K_3) = \lfloor n^2/4 \rfloor$ for all $n \geq 1$. Moreover, the only n-vertex K_3-free graphs with $\lfloor n^2/4 \rfloor$ edges are those which are isomorphic to the balanced complete bipartite graph $K_{\lfloor n/2 \rfloor, \lceil n/2 \rceil}$.*

There are many proofs for Mantel's Theorem (the textbook "Proofs from the Book" contains 7 proofs, and there are many more than just these!). We will content ourselves with only a single proof here, though we sketch out a few more in the exercises.

Proof. One reasonable approach to consider when given a problem like this is to try and prove things by induction on n, which is indeed what we shall ultimately do, though we will have to be a little careful with the details.

Indeed, consider the following naive approach using induction: let G be an n-vertex K_3-free graph and v an arbitrary vertex of G. Inductively we know that $e(G - v) \leq \lfloor (n-1)^2/4 \rfloor$, and hence $e(G) \leq \lfloor (n-1)^2/4 \rfloor + \deg(v)$. Unfortunately this bound is not good enough: if, say $G = K_{1,n-1}$ and v were the center of the star then this would give a bound of $\lfloor (n-1)^2/4 \rfloor + n-1$, which is too large. One can try and be smarter by picking v to be a vertex of minimum degree, but we do not know if this is enough to prove the result. To deal with this, we will prove the result by removing *two* vertices at a time from G rather than just one.

To this end, observe that the result is true for $n = 1, 2$. Assume we have proven the result up to some value $n \geq 3$ and let G be an n-vertex triangle-free graph. If $e(G) = 0$ then we are done, so we can assume G has an edge xy. By induction, we know that $e(G-x-y) \le \lfloor (n-2)^2/4 \rfloor = \lfloor n^2/4 \rfloor - n + 1$, and hence that

$$e(G) = e(G-x-y) + \deg(x) + \deg(y) - 1 \le \lfloor n^2/4 \rfloor + \deg(x) + \deg(y) - n,$$

where the -1 in the first equality comes from the fact that $xy \in E(G)$ and hence is counted by both $\deg(x)$ and $\deg(y)$. Finally, because G is triangle-free (which is a fact we must use somewhere in our argument), we must have $N(x) \cap N(y) = \emptyset$, as any common neighbor z would form a triangle with the edge xy. We conclude then that

$$\deg(x) + \deg(y) = |N(x)| + |N(y)| = |N(x) \cup N(y)| \le n,$$

which combined with the bound above give the desired bound.

To prove the equality case, again one can show this holds for $n = 1, 2$. Inductively then, the only way for the bound $e(G-x-y) \le \lfloor (n-2)^2/4 \rfloor$ to be tight is if $G-x-y = K_{\lfloor n/2 \rfloor -1, \lceil n/2 \rceil -1}$, and similarly the only way the bound $|N(x) \cup N(y)| \le n$ can be tight is if every vertex of $G-x-y$ is adjacent to exactly one of x, y, which is only possible if G is $K_{\lfloor n/2 \rfloor, \lceil n/2 \rceil}$. $\square$

Similar to how the "correct" way to generalize our bound for C_4 in Theorem 1.1 was to consider complete bipartite graphs, it turns out that the "correct" way to generalize Mantel's Theorem is to consider larger cliques K_r. And indeed, just like the case of triangles, the Turán number for cliques in general can be solved exactly and has a unique extremal construction which is defined as follows.

Definition 3. Given integers $r, n \ge 1$, we define the *Turán graph* $T_{r-1}(n)$ to be the $(r-1)$-partite graph whose part sizes are as equal as possible, i.e. such that each part either has size $\lfloor n/(r-1) \rfloor$ or size $\lceil n/(r-1) \rceil$

Theorem 1.8 (Turán's Theorem). *For all integers $r \ge 2$ and $n \ge 1$, we have $\mathrm{ex}(n, K_r) = e(T_{r-1}(n))$. Moreover, the only n-vertex K_r-free graph with $e(T_{r-1}(n))$ edges are those which are isomorphic to $T_{r-1}(n)$.*

Again there are many different proofs of Turán's Theorem, and again we limit ourselves to just a single one here based on the following idea.

Mantra 5. If you think an extremal problem has a unique optimal construction, then try and prove this by "shifting" an arbitrary construction to look like the optimal construction.

For example, in the setting of Turán's Theorem we might want to try shifting an arbitrary K_r-free graph into a graph that, like the Turán graph $T_{r-1}(n)$, is complete $(r-1)$-partite. And indeed this is always possible to do.

Lemma 1.9 (Zykov Symmeterization). *For every K_r-free graph G, there exists a graph G' satisfying the following:*

- $V(G') = V(G)$,
- $\deg_{G'}(x) \ge \deg_G(x)$ *for all $x \in V(G)$, and*

- *G' is complete $(r-1)$-partite.*

In particular, these last two conditions imply that G' is a K_r-free graph with at least as many edges as G.

Emphasize somewhere how the high-level idea of the proof is to duplicate vertices of high degree while deleting certain vertices of low degree.

Proof. We prove the result by induction on r, the case $r = 2$ being trivial. Let $x \in V(G)$ be a vertex of maximum degree. Observe that $H := G[N(x)]$ must be K_{r-1}-free, as any K_{r-1} in H together with x would form a K_r. By induction we can find a complete $(r-2)$-partite graph H' satisfying the conditions of the lemma for H. Now define G' to be the graph formed by starting with H' and then adding every edge from $V(H') = N_G(x)$ to the remaining vertices $x \cup (V(G) \setminus N_G(x))$.

Observe that $V(G') = V(G)$ and that G' is complete $(r-1)$-partite (namely by considering the $r-2$ parts from H' together with the part $x \cup (V(G) \setminus N_G(x))$, so it remains to check the degree condition. If $y \notin V(H') = N_G(x)$ then

$$\deg_{G'}(y) = v(H') = \deg_G(x) \geq \deg_G(y),$$

with this last inequality using that x was chosen to be a vertex of maximum degree. If instead $y \in V(H')$ then

$$\deg_{G'}(y) = \deg_{H'}(y) + |V(G) \setminus N_G(x)| \geq \deg_H(y) + |N_G(y) \setminus N_G(x)| = \deg_G(y),$$

where the inequality used $\deg_{H'}(y) \geq \deg_H(y)$ by definition of H. □

We now use this result to prove Turán's Theorem, though for simplicity we omit the proof of uniqueness.

Proof of Turán's Theorem. Let G be an n-vertex K_r-free graph. By Zykov symmeterization, we know that there exists an n-vertex complete $(r-1)$-partite graph G' with at least as many edges as G, and it is a simple exercise to show that any such graph has at most as many edges as $T_{r-1}(n)$, proving the result. □

As a historical aside, Turán proved this result without being aware of Mantel's Theorem, and in this paper he went on to introduce the general problem of determining $\mathrm{ex}(n, F)$ for various graphs F, which is why the "Turán number" bears his name.

1.3 Forbidding Trees

We have now solved the Turán problem for the "densest" graphs K_r. We now turn to solving the problem for the "sparsest" graphs, namely that of forests and trees. The simplest case of this problem is that of stars, which is easy to solve exactly.

Proposition 1.10. *For all $r \geq 2$, we have $\mathrm{ex}(n, K_{1,r-1}) \leq \frac{r-2}{2}n$ with equality if and only if at least one of r or n is even.*

Proof. A graph G being $K_{1,r-1}$-free is the same as saying that G has maximum degree at most $r-2$. Thus, any n-vertex $K_{1,r-1}$-free graph satisfies

$$e(G) = \frac{1}{2}\sum \deg(x) \le \frac{1}{2}\sum r-2 = \frac{r-2}{2}n,$$

proving the upper bound. This upper bound is tight whenever there exists an n-vertex $(r-2)$-regular graph, which holds precisely if at least one of r or n is even. □

Note that in this example there are infinitely many extremal constructions, which is a significantly different phenomenon compared to what we saw when forbidding cliques.

We next turn to the problem of avoiding an arbitrary tree T, for which we might ideally like to generalize our argument for stars. Unfortunately unlike in this case we can not say that an arbitrary T-free graph has small maximum degree, but we can prove the slightly weaker statement that such a graph has small minimum degree.

Lemma 1.11. *If T is a tree with r vertices and if G is a graph with minimum degree at least $r-1$, then G contains a copy of T.*

Note that the bound of $r-1$ is best possible, as can be seen by considering graphs G which are disjoint unions of copies of K_{r-1}. We present two essentially equivalent proofs of this result, the first of which is a little vaguer but requires less knowledge of trees while the second is a bit more explicit/algorithmic.

First Proof. We prove the result by induction on r, the case $r=2$ being trivial. Assume we have the proven the result up to some $r \ge 3$ and let T be an arbitrary r-vertex tree.

Because T is a tree, there exists some leaf x with some vertex y its unique neighbor. Because G has minimum degree at least $r-1 \ge r-2$, we inductively can assume that G has a copy of $T' = T - x$. Now the vertex playing the role of y in this copy of T' has at least $r-1$ neighbors, of which at most $r-2$ of them lie in this copy of T'. In particular, there exists at least one neighbor which is not in T', and taking this together with the copy of T' gives a copy of T gives the desired result. □

Second Proof. We build up our copy of T algorithmically "vertex by vertex." To do this we require the fact that for every r-vertex tree, there exists an ordering of the vertices $v_1, \ldots, v_r$ such that for all $2 \le i \le r$ there exists an integer $j_i < i$ such that $N_T(v_i) \cap \{v_1, \ldots, v_{i-1}\} = \{v_{j_i}\}$.

Let y_1 be an arbitrary vertex of G. Iteratively given that we have chosen vertices $y_1, \ldots, y_{i-1}$ in G for some $i \le r$, we choose y_i to be an arbitrary vertex in $N_G(y_{j_i})$ which is not in the set $\{y_1, \ldots, y_{i-1}\} \setminus \{y_{j_i}\}$. Note that the number of such vertices is at least $r-1-(i-2) \ge 1$, so there does indeed exist a valid choice for y_i, and as such this algorithm will successfully terminate. With this, it is not difficult to see that the y_i vertices form a copy of T, giving the result. □

The result above gives a tight bound on the minimum degree needed to contain a copy of T, but we ultimately want a bound on $\mathrm{ex}(n, T)$, i.e. on the *average* degree needed to find a copy of T. Fortunately, there is a general result which allows us to translate between the concept of minimum degrees and average degrees.

Proposition 1.12. *If G is a graph of average degree at least d, then there exists a non-empty subgraph $G' \subseteq G$ with minimum degree at least $d/2$ and average degree at least d.*

For most applications of this result we will only need the conclusion that G' has large minimum degree, but sometimes it is useful to also have this additional average degree condition (see for example Theorem 2.8). Again we offer two essentially equivalent proofs of this result, both of which implicitly use that the average degree by definition is

$$v(G)^{-1}\sum \deg(x) = \frac{2e(G)}{v(G)}.$$

First Proof. Assume the result is false for a given d and graph G, and choose such a counterexample with $v(G)$ as small as possible. If $\delta(G) \geq d/2$ then taking $G' = G$ gives the desired subgraph, a contradiction. As such, we can assume that G contains a vertex x with $\deg(x) < d/2$. In this case, the graph $G-x$ has a smaller number of vertices and average degree

$$\frac{2e(G-x)}{v(G-x)} = \frac{2e(G) - 2\deg(x)}{v(G)-1} \geq \frac{2e(G)-d}{v(G)-1} \geq d,$$

with this last step using that $2e(G) \geq dv(G)$ by hypothesis. Since $G-x$ is a graph with fewer vertices than G and with average degree d, our choice of G having $v(G)$ as small as possible implies that there exists $G' \subseteq G - x \subseteq G$ satisfying the properties of the statement, giving another contradiction. □

Second Proof. The key idea of the argument is to start with $G' = G$ and then iteratively remove vertices of low degree, i.e. as long as G' contains a vertex of degree less than $d/2$ then we remove this vertex and we repeat this until no such vertices exist. Note that the total number of edges that we remove in this process is certainly less than

$$(d/2)\cdot v(G) \leq e(G),$$

with this inequality being equivalent to saying that G has average degree at least $d/2$. As such, the resulting graph G' has at least one edge and has minimum degree at least $d/2$ by construction. One can similarly check that it has average degree at least d, proving the result. □

This in total lets us prove the following.

Theorem 1.13. *For any r-vertex tree T, we have*

$$\frac{r-2}{2}n - O_r(1) \leq \mathrm{ex}(n,T) \leq (r-2)n$$

Proof. For the lower bound we take the disjoint union of copies of K_{r-1}, which is certainly T-free and which has the stated number of edges.

For the lower bound, assume that there exists an n-vertex T-free graph G with $e(G) > (r-2)n$, i.e. with average degree more than $2(r-2)$. By Proposition 1.12 there exists a subgraph G' of G with minimum degree more than $r-2$, i.e. with $\delta(G') \geq r-1$. By Lemma 1.11 $G' \subseteq G$ contains a copy of T, a contradiction. □

While Theorem 1.13 solves the Turán problem for trees up to a factor of 2, one can ask if one can give an even more precise answer. In particular, given that the lower bound of Theorem 1.13 is the truth for the case of stars, it is natural to believe this should be the answer in general.

Conjecture 1.14 (Erdős-Sós)**.** *Every r-vertex tree T satisfies $\mathrm{ex}(n,T) \le \frac{r-2}{2}n$.*

There are a number of special cases for which the Erdős-Sós Conjecture is known to be true (such as for paths; see Theorem 2.8), but overall the problem of improving the small gap from Theorem 1.13 for all T seems difficult to do

1.4 An Aside: Degenerate vs Non-Degenerate Turán Problems

At this point we've studied $\mathrm{ex}(n,F)$ for a lot of classes of graphs F, but still we've said almost nothing about graphs in general. Part of the issue is that the Turán number can behave in very different ways depending on the structure of the graph F, in the following sense.

Proposition 1.15. *Let F be a graph.*

- *If F is non-bipartite, then $\mathrm{ex}(n,F) = \Theta(n^2)$.*
- *If F is bipartite, then $\mathrm{ex}(n,F) = O(n^{2-1/v(F)})$.*

Proof. For any graph we have $\mathrm{ex}(n,F) \le e(K_n) = \binom{n}{2} = O(n^2)$. If F is further non-bipartite, then the balanced complete bipartite graph $K_{\lfloor n/2\rfloor,\lceil n/2\rceil}$ is F-free and shows that $\mathrm{ex}(n,F) \ge \lfloor n^2/4\rfloor = \Omega(n^2)$, proving the first part. For the second part, because F is bipartite, we have $F \subseteq K_{v(F),v(F)}$, and hence by Kővári-Sós-Turán,

$$\mathrm{ex}(n,F) \le \mathrm{ex}(n, K_{v(F),v(F)}) = O(n^{2-1/v(T)}).$$

□

This observation divides the study of Turán number into two distinct cases: the *non-degenerate* case which studies non-bipartite F (i.e. those graphs with $\mathrm{ex}(n,F) = \Theta(n^2)$), and the *degenerate* case which studies bipartite F (i.e. those graphs with $\mathrm{ex}(n,F) = o(n^2)$). In what follows we very briefly survey some each of these cases. Some of these results require some real machinery to prove, and as such will be deferred until much later in the text.

The Non-Degenerate Case. For non-bipartite graphs F, the most important theorem is undoubtedly the following.

Theorem 1.16 (Erdős-Stone-Simonovits)**.** *For any graph F with at least one edge, we have*

$$\mathrm{ex}(n,F) = \left(\frac{\chi(F)-2}{\chi(F)-1} + o(1)\right)\binom{n}{2}.$$

In particular, this result determines the asymptotic value of $\mathrm{ex}(n,F)$ for *any* non-bipartite[3] graph F. The lower bound for this is rather easy: the Turán graph $T_{\chi(F)-1}(n)$ has chromatic

[3]If F is bipartite the theorem simply says $\mathrm{ex}(n,F) = o(n^2)$, which follows from Kővári-Sós-Turán.

number $\chi(F) - 1$ and hence is F-free and has about $\frac{\chi(F)-2}{\chi(F)-1}\binom{n}{2}$ edges. The upper bound is somewhat difficult to prove, and we will defer this until we have the powerful machinery of the regularity lemma at our disposal.

Because of the power of the Erdős-Stone-Simonovits theorem, the non-degenerate case of the Turán problem is often considered to be largely solved. That being said, one can ask for your favorite non-bipartite graph F if one can prove sharper (possibly even exact) bounds on $\mathrm{ex}(n, F)$ or to determine the full set of optimal extremal constructions. There are a number of results in this direction, with perhaps the most useful being the following result of Simonovits.

Theorem 1.17. *Let F be a graph which is "edge-critical", meaning it contains an edge e with $\chi(F - e) < \chi(F)$. Then $\mathrm{ex}(n, F) = e(T_{\chi(F)-1}(n))$ for all n sufficiently large, and moreover the unique extremal construction for n sufficiently large is $T_{\chi(F)-1}(n)$.*

We emphasize that the assumption of n being sufficiently large is necessary in general. Indeed, we always have $\mathrm{ex}(n, F) = \binom{n}{2}$ whenever $n < v(F)$, and for small n this will typically be better than the bound given in Theorem 1.17.

Finally, we note that while the Erdős-Stone-Simonovits Theorem largely solves the case of non-degenerate Turán problems for graphs, the analogous problem for *hypergrpahs* remains very wide open. We'll touch on this a bit more somewhere later.

The Degenerate Case. While non-degenerate Turán problems for graphs are largely solved, nothing could be farther from the case for degenerate Turán problems. Indeed, even determining the order of magnitude of relatively simple bipartite graphs remain open despite decades of study. We already mentioned that for complete bipartite graphs that $\mathrm{ex}(n, K_{s,s})$ remains open for all $s \geq 4$. Similarly, for even cycles (which are perhaps the next most natural class of bipartite graphs to study) our knowledge is summarized as follows.

Theorem 1.18. *For all $\ell \geq 2$, we have $\mathrm{ex}(n, C_{2\ell}) = O(n^{1+1/\ell})$. Moreover, this is known to be best possible whenever $\ell = 2, 3$, or 5.*

That is, we know the Turán number for C_4, C_6, and C_{10}, but frustratingly not for C_8! This is roughly because there happen to exists a class of very particular algebraic objects which just so happen to solve these cases and no others. Another frustrating problem is that of the 3-dimensional hypercube graph Q_3, which can be viewed as the "skeleton" of a usual cube. Determining $\mathrm{ex}(n, Q_3)$ was one of the original problems that Turán raised back in his 1941 paper on the topic, but to date only the following bounds are known.

Theorem 1.19. *We have $\mathrm{ex}(n, Q_3) = O(n^{8/3})$ and $\mathrm{ex}(n, Q_8) = \Omega(n^{3/2})$.*

The lower bound comes simply by considering an extremal C_4-free graph. The upper bound is based on a "supersaturation" argument of Erdős and Simonovits from 1969.

Much more can be said about what we do not know about Turán numbers of bipartite graphs, see Survey.

1.5 Exercises

1. Verify that the graphs G_q, G_q^* defined in the first subsection are C_4-free and that $v(G_q^*) = q^2+q+1$ and $e(G_q^*) = \frac{1}{2}(q+1)(q^2+q+1)$ [1+].

2. Prove the Kővári-Sós-Turán Theorem, Theorem 1.4 [1+].

3. Given integers $m,n,s,t \geq 1$, define the *Zarankiewicz number*[4] $z(m,n;s,t)$ to be the maximum number of edges in a bipartite graph G with parts U,V satisfying $|U|=m, |V|=n$, and that G no copy of $K_{s,t}$ with the part of size s in U and the part of size t in V.

   (a) Prove that
   $$z(m,n;s,t) \leq (t-1)^{1/s} mn^{1-1/s} + (s-1)n.$$
   (Hint: if you're struggling with this, try solving the previous problem first) [2].

   (b) Prove that if G is an n-vertex bipartite C_4-free graph then $e(G) \leq 2^{-3/2}n^{3/2} + o(n^{3/2})$, i.e. the lower bound we got for $\mathrm{ex}(n,C_4)$ using G_q was best possible in the setting of bipartite graphs [2-].

   (c) Prove that for all s,t there exists a constant $C>0$ such that if G is an n-vertex $K_{s,t}$-free graph, then the number of edges $xy \in E(G)$ with $\deg(x) \geq Cn^{1-1/s}$ is at most $O(n)$. Find an example of a graph which has $\Theta(n)$ edges of this form (Hint: the intended proof I have in mind works with $C \approx (s+t-1)^{1/s}$) [2].

   (d) Use (a) with $s=t=2$ to give a generalization of Corollary 1.2 [1].

4. The Turán problem involves graphs with 0 copies of a given graph F (where here by a *copy* we mean a subgraph isomorphic to F). What about graphs with more copies?

   (a) Prove that if G is an n-vertex graph then G contains at least $e(G) - \mathrm{ex}(n,F)$ copies of F for any graph F [1].

   (b) Prove that if G is an n-vertex graph with $e(G) \geq 100n^{3/2}$ then G contains at least $\Omega(n^{-6}e(G)^4)$ copies of C_4 (The number 100 does not matter in case you'd rather prove this result with a different constant) [2].

   Note that the number of copies guaranteed in (b) is far more than the naive bound given by (a). This sort of phenomenon of graphs with $e(G)$ just above $\mathrm{ex}(n,F)$ having a surprisingly large jump in the number of copies of F is known as *supersaturation*.

   (c) Prove that for all m with $100n^{3/2} \leq m \leq \binom{n}{2}$ that there exists an n-vertex graph G with $e(G) = \Theta(m)$ and with $\Theta(n^{-6}m^4)$ copies of C_4 (Hint: consider something random) [2+].

5. Prove that $\mathrm{ex}(n, K_{3,3}) = \Omega(n^{5/3})$ [3].

[4]Some texts define $z(m,n;s,t)$ with respect to G which are $K_{s,t}$-free rather than simply avoiding things on one side like we have here.

6. Prove that $\mathrm{ex}(n, K_{s,t}) = \Omega(n^{2-1/s})$ for all t sufficiently large in terms of s [3+].

* * *

7. Determine $\mathrm{ex}(n, F)$ for all graphs F with $2 \leq v(F) \leq 3$ other than $F = K_3$. Why did I leave out the case $v(F) = 1$? [1].

8. Verify that if G' is an n-vertex complete $(r-1)$-partite graph then $e(G') \leq e(T_{r-1}(n))$ [1+].

9. Here we sketch a few alternative proofs of Mantel's Theorem and Turán's Theorem.

   (a) Observe that if G is a triangle-free graph, then $\deg(x) + \deg(y) \leq v(G)$ for all $xy \in E(G)$. Use this to prove Mantel's Theorem (which is in fact the original way Mantel proved his result) [2].

   (b) Generalize our inductive proof of Mantel's Theorem to give an alternative proof of Turán's Theorem (which is in fact the original way that Turán proved his result). For simplicity you can choose to prove only that

   $$\mathrm{ex}(n, K_r) \leq \left(1 - \frac{1}{r-1}\right)\frac{n^2}{2},$$

   which one can check is equivalent to proving the upper bound of Turán's Theorem [2].

10. Let F denote the unique 4-vertex graph with 5 edges (i.e. the graph consisting of two triangles sharing an edge). Prove (without using Theorem 1.17) that $\mathrm{ex}(n, F) = \lfloor n^2/2 \rfloor$ for all $n \geq 4$ [2].

11. If F denotes the "bowtie" graph consisting of two triangles sharing a vertex, show that $\mathrm{ex}(n, F) = \lfloor n^2/2 \rfloor + 1$ for all $n \geq 6$ [3-].

* * *

12. Determine $\mathrm{ex}(n, P_4)$ exactly for all n (Hint: characterize all connected P_4-free graphs) [2].

13. Prove that for every integer $s \geq 1$ and real $\varepsilon > 0$, there exists a graph with average degree at least $2s - \varepsilon$ which contains no non-empty subgraph with minimum degree greater than $s+1$; that is, the $d/2$ in Proposition 1.12 is essentially best possible [2].

* * *

14. One can consider Turán problems which avoids more than just a single graph at a time. To this end, given a set of graphs $\mathcal{F}$, we say that a graph G is $\mathcal{F}$-free if G is F-free for all $F \in \mathcal{F}$.

    Prove that for all $\ell \geq 2$ we have $\mathrm{ex}(n, \{C_3, C_4, \ldots, C_{2\ell}\}) = O(n^{1+1/\ell})$ (Hint: first prove the result under the additional assumption that every vertex of G has degree at least $n^{1/\ell}$) [2].

15. Prove that if F is a graph with $\mathrm{ex}(n, F) = \Omega(n)$ and if F' is a graph obtained from F by adding a new vertex x and making it adjacent to a vertex $y \in V(F)$, then $\mathrm{ex}(n, F') = \Theta(\mathrm{ex}(n, F))$. In other words, to determine the order of magnitude of $\mathrm{ex}(n, F)$ for all graphs F, it suffices to do so for all graphs with minimum degree at least 2 [2].

16. (Füredi) Prove that if F is a bipartite graph where every vertex on one side of the bipartition has degree at most r, then $\mathrm{ex}(n, F) = O(n^{2-1/r})$. Show that this bound is best possible for all r [3+].

2 Spanning Subgraphs and Dirac Problems

Up to this point we have considered the Turán number $\mathrm{ex}(n, F)$ where we think of F as a fixed graph and n as tending towards infinity, but this is not the only regime that could be considered. For example, $\mathrm{ex}(n, C_n)$ asks for the maximum number of edges that an n-vertex graph can have without containing a Hamiltonian cycle. More generally, we might consider $\mathrm{ex}(n, F_n)$ where F_n is some sequence of spanning subgraphs of K_n.

Unfortunately the Turán problem for spanning subgraph tends not to be very interesting. For example, one can show $\mathrm{ex}(n, C_n) \geq \binom{n-1}{2} + 1$ by taking G to be a clique on $n-1$ vertices together with a single vertex of degree 1, and it is not too difficult to show that this somewhat silly construction is best possible. More generally, $\mathrm{ex}(n, F_n)$ tends to be ludicrously large for a number of natural choices of F_n simply by considering graphs G which have a single vertex of small degree. This leads us to another mantra.

Mantra 6. If an extremal problem has a known or boring optimal construction, try modifying or adding extra restrictions to the problem in such a way that any solution to this new problem must be "far" from the known/boring construction.

In particular, our current construction for $\mathrm{ex}(n, C_n)$ is boring because we can trivially make constructions by using vertices of very small degrees. So what if we instead forced our constructions to all have large minimum degree? This leads to the following broad type of problem.

Definition 4. Given a graph F, we define[5] the *Dirac number* $\delta^*(F)$ to be the smallest number δ^* such that any $v(F)$-vertex graph G with $\delta(G) \geq \delta^*$ has a copy of F as a spanning subgraph.

Note that we have already seen some problems somewhat similar to δ^* when we were working on Turán numbers for trees via Lemma 1.11. We will see another application of min degree results to Turán problems with Theorem 2.8.

2.1 Hamiltonian Cycles

Recall that a graph G is Hamiltonian if it contains a cycle passing through all of its vertices. Historically, the first study of Dirac numbers came from Dirac who determined $\delta^*(C_n)$, i.e. the smallest minimum degree of an n-vertex graph G which guarantees that G is Hamiltonian.

To start our investigation, let us try to think of some graphs with large minimum degree which do not have a Hamiltonian cycle. One immediate way to tell that a graph does not have a Hamilotonian cycle is if the graph is disconnected. In particular, if we consider G to be the n-vertex graph which is the disjoint union of $K_{\lceil n/2 \rceil}$ and $K_{\lfloor n/2 \rfloor}$, then this is a graph with no Hamiltonian cycle and with minimum degree $\lfloor n/2 \rfloor - 1$, showing that we must have $\delta^*(C_n) \geq \lfloor n/2 \rfloor$. While perhaps not as obvious, there exists another construction that gives a very similar bound which one might discover by looking at the cases of small n, for example. Specifically, any graph of the form $K_{m,n-m}$ with $m < \lceil n/2 \rceil$ will fail to be Hamiltonian. Indeed,

[5]This name is completely made up since we are not aware of any standard name for this parameter in the literature.

if n is odd this is immediate because $K_{m,n-m}$ is bipartite and hence can not contain C_n. If n is even then any Hamilton cycle in such a graph must have exactly $n/2 = \lceil n/2 \rceil$ of its vertices lying in each part of $K_{m,n-m}$, which is impossible to do under the condition $m < \lceil n/2 \rceil$. This construction thus implies that $\delta^*(C_n) \geq \lceil n/2 \rceil$, which matches the bound in the previous construction if n is even and does a little better if n is odd. In total it turns out that this bound is indeed the correct one.

Theorem 2.1 (Dirac's Theorem). *Every n-vertex graph G with $\delta(G) \geq n/2$ contains a Hamiltonian cycle.*

The reader should double check that this, together with our constructions from above, is equivalent to saying that $\delta^*(C_n) = \lceil n/2 \rceil$. Before we get on with the proof, let us make the meta-observation that for n even there are two extremal constructions for Dirac's Theorem (the disjoint union of two equally sized cliques, and a slightly unbalanced complete bipartite graph). This is non-ideal due to the following

Mantra 7. Extremal problems tend to be harder if they have more than one extremal constructions, especially if these constructions look very different from each other.

Indeed, part of the ease of proving Turán's Theorem is that there is only one possible extremal construction, which means we can hope to do arguments like Zykov symmeterization which move us closer to this unique extremal example. However, this approach as well as many others fail when there are multiple different looking extremal examples because whatever argument we make must simultaneously be optimal for all of our possible constructions.

To partially deal with this issue, we will utilize another mantra.

Mantra 8. If during a proof you assume that there exists some counterexample to your statement, it is sometimes useful to assume this counterexample is "extremal" in some sense.

We will see a concrete example of this in our following proof of Dirac's Theorem, which is originally due to Posa maybe.

Proof of Dirac's Theorem. Assume for some integer n that there exists a counterexample G and, crucially, choose such a counterexample with as many edges as possible. Intuitively by choosing a graph with more edges should make it easier for us to construct a Hamiltonian cycle, giving the desired contradiction. In particular, this assumption gives us the following key fact.

Claim 2.2. *The graph G contains a Hamiltonian path $x_1 \cdots x_n$.*

Proof. This is trivial if $G = K_n$, so assume this is not the case, i.e. that there exists some non-edge $xy \notin E(G)$. Because $G + xy$ is an n-vertex graph with $\delta(G + xy) \geq \delta(G) \geq n/2$ and with strictly more edges than G, it must be that $G + xy$ contains a Hamiltonian cycle C by assumption of G being a counterexample with the maximum number of edges. The subgraph $C - xy$ then must be a Hamiltonian path. □

The other key observation we will need is the following.

Claim 2.3. *If there exists an integer $2 \le i \le n$ such that $x_i \sim x_1$ and $x_{i-1} \sim x_n$, then G is Hamiltonian.*

Proof. Consider the following sequence of vertices:

$$P = (x_1, x_i, x_{i+1}, \ldots, x_{n-1}, x_n, x_{i-1}, x_{i-2}, \ldots, x_2).$$

It is not difficult to see that P is a Hamiltonian path (i.e. every vertex appears exactly once and consecutive vertices are adjacent) with its first and last vertices being adjacent to each other. Therefore this defines a Hamiltonian cycle in G, proving the claim. $\square$

As an aside, the idea in this claim of "rotating" the Hamiltonian path we started with into a new one P is a common idea known as a Pósa rotation.

Back to our problem at hand, we want to show that an index i as in the claim exists. To this end, define

$$X_1 = \{i : x_i \sim x_1\},$$
$$X_n = \{i : x_{i-1} \sim x_1\}.$$

By the claim above and our assumption that G is Hamiltonian, we can assume that X_1, X_n are disjoint subsets of $\{2, \ldots, n\}$. This implies that

$$n - 1 \ge |X_1 \cup X_n| = |X_1| + |X_n| = \deg(x_1) + \deg(x_n) \ge n,$$

a contradiction. $\square$

Even though Dirac's Theorem is tight, it is still possible to ask for strengthenings of this result as follows.

Mantra 9. After proving a theorem, check to see where you use the hypothesis of your theorem and if these can be relaxed in any way.

For example, the only place where we really used $\delta(G) \ge n/2$ in our proof of Dirac's Theorem was to show that $\deg(x_1)+\deg(x_n) \ge n$. A moments thought then shows that our proof actually implies the following stronger result.

Theorem 2.4 (Ore's Theorem). *If G is an n-vertex graph such that for every non-edge $xy \notin E(G)$ we have $\deg(x) + \deg(y) \ge n$, then G is Hamiltonian.*

In fact, our proof has much more flexibility that can be exploited to prove a number of other extensions. We state another one here and leave its proof as an exercise to the reader.

Theorem 2.5 (Pósa's Theorem). *If G is an n-vertex graph such that for all integers $k < n/2$,*

$$|\{x \in V(G) : \deg(x) \le k\}| < k,$$

then G is Hamiltonian.

These extensions of Dirac's Theorem, in addition to being nice on their own, also have various applications to them, such as the following.

Theorem 2.6. *If G is an n-vertex graph with $\delta(G) \geq \frac{n+1}{2}$, then for every edge $xy \in E(G)$ there exists a Hamiltonian cycle in G which uses the edge xy.*

Proof. Let $xy \in E(G)$ be an arbitrary edge, and consider a new graph G' obtained by adding a new vertex v which is adjacent to only x, y. This $(n+1)$-vertex graph G' satisfies the conditions of Pósa's Theorem (it has only 1 vertex of degree at most 2, and every other vertex has degree at least $v(G')/2$), so G' contains a Hamiltonian cycle C. Note that this Hamiltonian cycle must contain the edges xv, vy since these are the only two neighbors of v. As such, the graph $C - v + xy$ is a Hamiltonian cycle in G using the edge xy, proving the result. □

2.2 Applications to Paths

Having just determined the optimal minimum degree needed to guarantee a graph contains a Hamiltonian cycle, it is natural to ask what conditions guarantee a Hamiltonian path. In fact, this turns out to be a consequence of Dirac's Theorem.

Theorem 2.7. *We have $\delta(P_n) = \lfloor n/2 \rfloor$. Equivalently, any n-vertex graph G with $\delta(G) \geq \frac{n-1}{2}$ contains a Hamiltonian path and this bound is best possible.*

Proof. The fact that this bound is best possible follows by considering G to be the disjoint union of two cliques of sizes $\lfloor n/2 \rfloor, \lceil n/2 \rceil$.

Now let G be an n-vertex graph with $\delta(G) \geq \frac{n-1}{2}$ and consider a new graph G' obtained by adding a vertex v which is adjacent to every vertex of G. Then $\delta(G') \geq (n+1)/2 = v(G')/2$, so by Dirac's Theorem G' contains a Hamiltonian cycle C, and thus $C - v$ is a Hamiltonian path in G. □

The trick we used in the proof above lets us easily translate many of the results that we have for Hamiltonian cycles to that of Hamiltonian paths; see the exercises for more.

We can also use Dirac's Theorem to prove good bounds for Turán numbers of paths.

Theorem 2.8 (Erdős-Gallai)**.** *For all $r \geq 2$, we have $\mathrm{ex}(n, P_r) \leq \frac{r-2}{2} n$.*

Note that this bound is tight whenever $r - 1 | n$, as can be seen by considering G to be the disjoint union of K_{r-1}'s.

Proof. By prove the result by double induction on r and n. The result for all n is trivial when $r = 2$, so assume we have proven the result for all n up to some value r. This result in turn is trivial if $n \leq r - 1$, so we assume we have the proven the result up to some value $n \geq r$. With this in mind, let G be an extremal n-vertex P_r-free graph and assume for contradiction that $e(G) > \frac{r-2}{2} n$.

Because our extremal example looks like a disjoint union of K_{r-1}'s, a perhaps reasonable thing to try and prove is the following.

Claim 2.9. *The graph G contains a cycle C with $r - 1$ vertices.*

Proof. By Proposition 1.12, there exists a subgraph $G' \subseteq G$ with minimum degree at least $\frac{r-1}{2}$ (i.e. strictly more than $\frac{r-2}{2}$) and average degree strictly more than $r-2$. By induction on r and the fact that G' has average degree more than $r-2$, we conclude that G' must contain a path $x_1 \cdots x_{r-1}$.

Now all of the neighbors for x_1, x_{r-1} must lie within $\{x_1, \ldots, x_{r-1}\}$, as otherwise $G' \subseteq G$ would contain a path on r vertices. Because $\deg_{G'}(x_1), \deg_{G'}(x_{r-1}) \geq \frac{r-1}{2}$, the exact same argument that we used in the proof of Dirac's Theorem implies that there exists a cycle C using all of the vertices in $\{x_1, \ldots, x_{r-1}\}$. □

Observe that every vertex in C can only be adjacent to other vertices of C, as one could use any additional neighbor together with C to construct a P_r in G. As such, the number of edges incident to the vertices of C is at most $\binom{r-1}{2}$, and as such the graph $G - V(C)$ is a smaller order graph which has

$$e(G - V(C)) > \frac{r-2}{2}n - \binom{r-1}{2} = \frac{r-2}{2}(n-r+1),$$

and since $G - V(C)$ has $n-r+1$ vertices, we conclude by induction on n that $G - V(C)$ has a P_r, giving the result. □

2.3 Clique Factors

Perhaps after Hamiltonian cycles and paths, the next most natural spanning structure to consider is that of a perfect matching, i.e. a disjoint union of K_2's which cover every vertex of the graph exactly once. Note that perfect matchings can only exist if the number of vertices in our graph is even.

While a natural problem to consider, perfect matchings will turn out to not be very interesting to study for two reasons. First, any graph with an even number of vertices and a Hamiltonian cycle (or path) contains a perfect matching, so by Dirac's Theorem we know that $\delta(G) \geq n/2$ is enough to guarantee a perfect matching, and this is best possible by considering $K_{n/2-1,n/2+1}$. Second, one can in fact characterize *exactly* when a given graph has a perfect matching as we shall see in later section, so just proving a sufficient condition is not so interesting.

While the exact problem of determining minimum degree conditions for perfect matchings is not so exciting, there are generalizations of perfect matchings for which this is very interesting. To this end, we say that a K_r*-matching* in a graph G is a subgraph of G which is the disjoint union of copies of K_r, and we say that G has a K_r*-factor* if G has a K_r-matching which contains every vertex of G exactly once. Note that G can only hope to have a K_r-factor if $r|n$.

Theorem 2.10 (Hajnal-Szemerédi Theorem Version I)**.** *If G is an n-vertex graph with $r|n$ and $\delta(G) \geq (r-1)n/r$, then G contains a K_r-factor.*

The Hajnal-Szemerédi Theorem is a deep result with a number of applications, see for example coloring chapter. The original proof of this result was quite difficult. There does exist a quite short proof due to Kierstead and Kostochka, but it is a little too dense to present here I think that's the case; double check. Rather than spending time on proving this in full, we will instead sketch out how to prove a somewhat weaker result.

Proposition 2.11. *If G is an n-vertex graph with $r|n$ and $\delta(G) \geq (r-1)n/r$, then G has a K_r-matching which contains all but at most $(r-1)^2 r$ vertices of G.*

Sketch of Proof. The rough idea is to consider a largest K_r-matching in G and argue that it has at least this size. However, to make the argument work we need to assume something slightly stronger about our matching.

To this end, let $S_1, \ldots, S_{n/r}$ be a partition of $V(G)$ into sets of size r such that $G[S_i]$ contains K_r for as many i as possible, and conditional on this, we choose this partition so that $G[S_i]$ contains a K_{r-1} for as many i as possible, and so on. Let $C_i \subseteq S_i$ denote a largest clique in $G[S_i]$ and assume for contradiction that $G[C_i] \neq K_r$ for at least $(r-1)^2+1$ values of i. By the Pigeonhole principle, this implies there is some $\ell \in [r-1]$ such that $|C_i| = \ell$ for at least r values of i, say for all $i \in [r]$ without loss of generality. Let $N(C_i)$ denote the set of common neighbors of C_i, i.e. the vertices adjacent to every vertex of C_i.

Claim 2.12. *We have $|N(C_i)| \geq (r-\ell)n/r$ and $N(C_i) \cap C_j = \emptyset$ for all $i, j \in [r]$.*

Proof. The lower bound $|N(C_i)| \geq (r-\ell)n/r$ follows from the fact that each of the ℓ vertices of C_i have minimum degree at least $(r-1)n/r$, i.e. are non-adjacent to at most n/r vertices. For the second part, assume for contradiction that there exists some $v \in N(C_i) \cap C_j$ and let $w \in S_j \setminus C_j$ be arbitrary (which exists since $|C_j| < r = |S_j|$). In this case, we could change our partition by replacing S_i, S_j with $S_i \cup \{v\} \setminus \{w\}$ and $S_j \setminus \{v\} \cup \{w\}$, which would increase the number of sets in the partition which contain a $K_{\ell+1}$ while not decreasing the number of sets containing any larger clique, contradicting how we chose our partition. We conclude that no such v exists. $\square$

In total this claim implies $\sum_{i=1}^r |N(C_i) \cap \bigcup_{j>r} C_j| \geq (r-\ell)n$, which by the Pigeonhole principle implies there is some $j > r$ such that

$$\sum_{i=1}^{r} |N(C_i) \cap C_j| \geq \left\lceil \frac{(r-\ell)n}{n/r - r} \right\rceil \geq r(r-\ell)+1.$$

Claim 2.13. *There exists some distinct $i', i'' \in [r]$ and disjoint $C_j', C_j'' \subseteq C_j$ of sizes 1 and $r-\ell$ such that $C_j' \subseteq N(C_{i'}) \cap C_j$ and $C_j'' \subseteq N(C_{i''}) \cap C_j$.*

Proof. By the inequality above and the Pigeonhole principle, there exists $i' \in [r]$ such that $|N(C_{i'}) \cap C_j| \geq r-\ell+1$, and since $|N(C_{i'}) \cap C_j| \leq r$ we have

$$\sum_{i \in [r] \setminus \{i'\}} |N(C_i) \cap C_j| \geq r(r-\ell-1)+1,$$

so again by the Pigeonhole principle there exists $i'' \neq i'$ such that $|N(C_{i''}) \cap C_j| \geq r-\ell$. Let $C_j'' \subseteq N(C_{i''}) \cap C_j$ be an arbitrary subset of size $r-\ell$ and let $C_j' \subseteq N(C_{i'}) \cap C_j$ be an arbitrary vertex disjoint from C_j', giving the result. $\square$

Let $w \in S_{i'} \setminus C_{i'}$ be arbitrary. If we consider modifying the partition by replacing $S_{i'}, S_{i''}, S_j$ (whose largest cliques have sizes ℓ, ℓ, r) with the r-sets $S_{i'} \cup C_j' \setminus \{w\}$, $C_{i''} \cup C_j''$, and $S_j \cup \{w\} \setminus$

$(C'_j \cup C''_j)$ (whose largest cliques have sizes at least $\ell+1, r, 1$), we see that this strictly increases the number of sets in our partition containing a $K_{\ell+1}$ while maintaining the sizes of all larger cliques, a contradiction to how we chose our partition. $\square$

We emphasize that for many Dirac-type problems it is relatively easy to find an "almost spanning" subgraph like we did here, but finding a genuinely spanning structure is often difficult. One general tool for doing this is the absorption method which we probably won't talk about, but we'll see what happens.

2.4 Exercises

1. Let's look at Turán numbers of spanning subgraphs.

   (a) Prove that $\mathrm{ex}(n, C_n) = \binom{n-1}{2} + 1$ [2].

   (b) Prove that $\mathrm{ex}(n, P_n) = \binom{n-1}{2}$ [2-].

2. We've seen that a minimum degree of about $n/2$ is the threshold for guaranteeing both a perfect matching and a Hamiltonian cycle. In the next few exercises we show that the behaviors for matchings and cycles differ greatly from one another when other sorts of degree conditions are imposed.

   (a) Prove that if G is an n-vertex graph with minimum degree $d \geq 2$, then G contains a cycle on at least $d+1$ vertices. Moreover, prove that for infinitely many n there exists an n-vertex graph which are $(d-1)$-regular and which have no cycle on at least $d+1$ vertices [1+].

   (b) Prove that if G is an n-vertex graph with minimum degree $d \geq 2$ and $d \leq n/2$, then G contains a matching on at least $2d$ vertices. Moreover, prove that this is best possible, i.e. that the result is false if we do not impose the condition $d \leq n/2$ and that there exist infinitely many $n \geq 2d$ with minimum degree $d-1$ with no matching on at least $2d$ vertices (Hint: the argument you use here can't be a direct analog of a proof of Dirac's Theorem since the previous part shows such an approach will fail for cycles) [2].

   (c) Prove that if G is an n-vertex graph with $d \geq 1$ and maximum degree Δ, then every maximal matching of G (i.e. every matching which is not a subset of any larger matching) has at least $\frac{d}{2\Delta}n$ vertices. Moreover, prove for all integers $1 \leq d \leq \Delta$ with d even that there exists a graph G with minimum degree d and maximum degree Δ which contains a matching on at most $\frac{d}{\Delta+1}v(G)$ vertices [2].

   (d) Prove that if G is an n-vertex graph with minimum degree $d \geq 1$ and maximum degree Δ, then G contains a matching on at least $\frac{d}{d+\Delta}n$ vertices. Moreover, prove for all integers $1 \leq d \leq \Delta$ that there exist graphs G with minimum degree d and maximum degree Δ such that no matching has size larger than $\frac{2d}{d+\Delta}v(G)$ (Hint: what would you need to assume about G for the same argument from (c) to give you the desired bound? Can you make this assumption here?) [2+].

3. The original proof of Dirac's theorem went as follows:

   (a) Define a *lollipop* to be a graph which consists of a cycle on vertices $v_1, \ldots, v_\ell$ together with a path on vertices $u_1, \ldots, u_t$ with $u_1 = v_1$. Given a graph G, consider its "largest" lollipop, i.e. the one which has ℓ as large as possible and conditional on this has t as large as possible.

   Prove that if such a largest lollipop has $\ell \geq 3$ and $t \geq 2$, then u_t is not adjacent to any two consecutive vertices in $v_1, \ldots, v_\ell$. Similarly prove that if $\ell \geq 3$ then u_t is not adjacent to any v_i vertex which is "close" to v_1. In particular, prove this is true for v_ℓ, v_2, then generalize this as much as you can (Hint: use the previous exercise) [2].

   (b) Conclude Dirac's Theorem [2]. .

4. Prove Pósa's Theorem [2].

5. Prove that if G is an n-vertex graph with $\delta(G) \geq (n+k)/2$ for some integer $k \geq 0$, then for any path $P \subseteq G$ on k edges there exists a Hamiltonian cycle of G which contains P as a subgraph (Hint: the trick we did before for $k = 1$ using Pósa's Theorem no longer works here, so you'll have to go back and modify our proof of Dirac's Theorem instead) [2].

6. Prove that if G is an n-vertex graph and $\delta(G) \geq n/2$, then for every edge of G there exists a Hamiltonian path of G containing this edge [1+].

3 Ramsey Theory

Turán's original motivation for the Turán problem came from another area of extremal combinatorics known as Ramsey theory. In a very abstract sense, Ramsey theory (which extends far beyond just that of graphs) aims to prove that every sufficiently large structure contains relatively simple and orderly substructures. The original problem, as well as the namesake of the theory, comes from the following foundational result of Ramsey[6] from **REF**.

Definition 5. A *red-blue edge coloring* of a graph G is a map $\chi : E(G) \to \{\text{red, blue}\}$. We say that such a coloring has a *monochromatic* K_n if there exists a subgraph of G isomorphic to K_n such that either every edge of the subgraph is colored red or if every edge of the subgraph is colored blue.

Theorem 3.1 (Ramsey's Theorem)**.** *For all integers $n \geq 1$, there exists a (finite) N such that every red-blue edge coloring of K_N contains a monochromatic K_n.*

Equivalently, this says that for all integers $n \geq 1$, there exists some (finite) N such that every N-vertex graph G either contains a clique of size n or an independent set of size n (as can be seen by coloring the edges of K_N red if they belong to G and blue otherwise). That is, large graphs can not simultaneously have arbitrarily large clique and independent numbers.

The original proof of Ramsey's Theorem does not give explicit bounds on the size of N, and the central problem ins Ramsey Theory is to get better bounds on this quantity.

Definition 6. We define the *(diagonal) Ramsey number* $R(n)$ to be the smallest integer N such that every red-blue edge coloring of K_N contains a monochromatic K_n.

There are many variants of this classical Ramsey number $R(n)$, several of which we will discuss below.

3.1 Classical Bounds

Let us start by working some small examples to give a little intuition for the problem in general. It is immediate that $R(1) = 1$ and $R(2) = 2$, so the first non-trivial case of the problem is to determine[7] $R(3)$.

Proposition 3.2. *We have $R(3) = 6$.*

Proof. The lower bound comes from giving a coloring of the edges of K_5 which does not contain a triangle. The unique way to do this is to take a $C_5 \subseteq K_5$ and color its edges red with the remaining edges (which also form a C_5) being colored blue. It is easy to check that such a coloring has no monochromatic triangle.

[6] Funnily enough Ramsey was not a combinatorialist but rather a logician, and to this day there is still a lot of work on Ramsey theoretic problems from the perspectives of both logic and combinatorics.

[7] Colloquially this result is known as the "party problem" due to the following interpretation of its statement: if there are 6 people at a party, then there exist 3 people there who either all know each other or who all do not know each other.

For the upper bound, consider an arbitrary red-blue coloring of the edges of K_6 and assume for contradiction that this did not contain a monochromatic triangle. Let u be an arbitrary vertex, and observe that u has 5 total edges incident to it each of which is given one of 2 colors, so by the pigeonhole principle at least 3 of the edges of u all have the same color, say without loss of generality that the edges uv_1, uv_2, uv_3 are all colored red. Now if any edge v_iv_j is colored red then u, v_i, v_j would form a red triangle, so we can assume that all of the edges v_iv_j are colored blue. But in this case v_1, v_2, v_3 forms a blue triangle, again yielding a contradiction. □

At its core, the reason that the upper bound proof worked is that if a red-blue coloring does not contain a monochromatic K_3, then the "red neighborhood" of any vertex u can not contain either a red K_2 nor a blue K_3. Building on this idea leads to the following definition.

Definition 7. Given integers m, n, we define $R(m, n)$ to be the smallest integer N such that if every edge of K_N is colored either red or blue, then there either exists a red K_m or a blue K_n.

For example, one can check that $R(2, 3) = 3$ which is implicitly what we used in our upper bound proof for $R(3)$. Generalizing this idea gives the following observation of Erdős and Szekeres.

Lemma 3.3 (Erdős-Szekeres)**.** *For all $m, n \geq 2$, we have*

$$R(m, n) \leq R(m-1, n) + R(m, n-1).$$

Proof. Let $N = R(m-1, n) + R(m, n-1)$ and assume for contradiction that there exists a red-blue edge coloring of K_N which does not contain a red K_m nor a blue K_n. Let u be an arbitrary vertex and let V_R denote the set of vertices v such that uv is colored red, and similarly define V_B. Note that $|V_R| + |V_B| = N - 1 = R(m-1, n) + R(m, n-1) - 1$, and that we must either have $|V_R| \geq R(m-1, n)$ or $|V_B| \geq R(m, n-1)$ (since otherwise $|V_R| + |V_B| \leq R(m-1, n) + R(m, n-1) - 2$).

First consider the case that $V_R \geq R(m-1, n)$. By definition of $R(m-1, n)$, the coloring on $K_N[V_R]$ must contain either a red K_{m-1} or a blue K_n. The latter case can not happen by assumption of our coloring, and if the former happens then this K_{m-1} together with u would form a red K_m, again giving a contradiction. A similar conclusion holds if $|V_B| \geq R(m, n-1)$, proving the result. □

Using this recurrence relation together with the boundary condition $R(1, n) = R(n, 1) = 1$ gives the following.

Theorem 3.4. *For all $m, n \geq 1$, we have*

$$R(m, n) \leq \binom{m+n-2}{m-1}.$$

Indeed, by induction on $m + n$ we have that

$$R(m, n) \leq R(m-1, n) + R(m, n-1) \leq \binom{m+n-3}{m-2} + \binom{m+n-3}{m-1} = \binom{m+n-2}{m-1},$$

with the last step being Pascal's identity. Finally, taking $m = n$ in this bound gives bounds for diagonal Ramsey numbers.

Corollary 3.5. *For all $n \geq 1$, we have*

$$R(n) \leq \binom{2n-2}{n-1} \leq 4^n.$$

Let us turn now to lower bounds, starting with an elementary bound.

Lemma 3.6. *We have $R(n) \geq (n-1)^2 + 1$.*

Proof. Color the edges of $R_{(n-1)^2}$ via breaking up the vertex sets in to $n-1$ parts $V_1, \ldots, V_{n-1}$ each of size $n-1$ and coloring all the edges within each part red and all the edges between two parts blue. It is easy to see that this avoids monochromatic copies of K_n. $\square$

Note that in this coloring that the blue edges form a copy of the Turán graph $T_{n-1}(n-1)$ and I think there's some connection here but I forget the details. It was believed for some time that $R(n)$ should grow polynomialy like in this lemma here, but Erdős disproved this in a very strong form.

Theorem 3.7. *We have*

$$R(n) \geq (1 + o(1)) \frac{n}{e\sqrt{2}} 2^{n/2}.$$

This is a strange bound and not one should necessarily expect to understand how to prove even if you work out the right proof idea. Indeed, our proof will utilize the following.

Mantra 10. First figure out how your proof works using an abstract set of parameters, then go back and choose whatever parameters you need in order for the arithmetic to go through

Let us see this in action.

Proof. To partially motivate the idea of the argument, we observe that it is very easy to show $R(n) \geq n+1$ for $n \geq 3$. Indeed, there are only two colorings of K_n which contain a monochromatic K_n, and as long as $n \geq 3$ we can find a coloring which avoids one of these two bad ones. To get our stated lower bound, we will similarly use an elementary counting argument to bound the number of "bad" colorings of K_N and then argue that if N is not too large then there are more total colorings than bad colorings, proving that there exists some coloring which is not bad.

From now on we fix an integer N which we will determine later once we see how the numbers work out. For each subset $S \subseteq [N]$ of size n, let B_S denote the set of edge colorings of K_N which have a monochromatic K_n on S. Because the total number of edge-colorings of K_N is $2^{\binom{N}{2}}$ and because a coloring avoids monochromatic K_n's if and only if it does not lie in any B_S set, we see that there exists an edge-coloring of K_N avoiding monochromatic K_n's if and only if

$$2^{\binom{N}{2}} > |\bigcup_{S \in \binom{[N]}{n}} B_S|.$$

It thus remains to show that this latter set is small. Using elementary arguments we have

$$|\bigcup_{S\in\binom{[N]}{n}} B_S| \le \sum_{S\in\binom{[N]}{n}} |B_S| = \binom{N}{n} 2^{1+\binom{N}{2}-\binom{n}{2}}$$

where this last step used that every coloring in B_S has 2 choices for how it can act on the edges of S (either all red or all blue) together with $2^{\binom{N}{2}-\binom{n}{2}}$ choices for the remaining edges. As such, we will succeed if

$$\binom{N}{n} 2^{1-\binom{n}{2}} < 1.$$

To get a handle on this, we use the well-known binominal inequality $\binom{m}{k} \le (em/k)^k$ to conclude that it suffices to have N such that

$$2\left(\frac{eN2^{(n-1)/2}}{n}\right) < 1,$$

and in particular the result holds provided $N < 2^{1/n} \cdot \frac{n}{e\sqrt{2}} 2^{-n/2}$, and picking such an N gives the desired bound. $\square$

This counting argument is all well and good, but we can give a more modern perspective by rewriting our proof in the language of probability.

Alternative Proof. Let N be an integer to be determined later and consider a uniform random red-blue edge coloring of K_N. Let X be the random variable which is equal to the number of monochromatic K_n's that are in the random coloring of K_N. Crucially, we observe that if $\mathbb{E}[X] < 1$, then $R(n) > N$. Indeed, because X is integer valued, the only way $\mathbb{E}[X] < 1$ is possible is if there exists some coloring of K_N such that $X = 0$, i.e. a coloring without any monochromatic copies of K_N.

To get a handle on $\mathbb{E}[X]$, for each $S \in \binom{[N]}{n}$ we let $\mathbb{1}_S$ denote the indicator random variable for $K_N[S]$ being monochromatic. That is, $\mathbb{1}_S$ is the random variable defined by having $\mathbb{1}_S = 1$ if $K_N[S]$ is monochromatic and $\mathbb{1}_S = 0$ otherwise. With this $X = \sum \mathbb{1}_S$, so by linearity of expectation we have

$$\mathbb{E}[X] = \sum \mathbb{E}[\mathbb{1}_S] = \sum \Pr[\mathbb{1}_S = 1] = \binom{N}{n} 2^{1-\binom{n}{2}},$$

as can be checked by a simple counting argument. Thus in total, we conclude $R(n) > N$ provided $\binom{N}{n}2^{1-\binom{n}{2}} < 1$, which as we showed in the previous version of the proof happens for $N = (1 + o(1))\frac{n}{e\sqrt{2}}2^{-n/2}$. $\square$

While both the counting argument and the probabilistic argument for [theorem] are effectively equivalent to each other, the prespective of "thinking probabilistically" has proven to be the more useful in general. Indeed, it is hard at this point not to find an important result in Ramsey theory where the lower bound (and sometimes even the upper bound) does not use some amount of ideas or techniques motivated by probability theory. Since we are not assuming the reader has any knowledge of probability we will not dwell on this point any further at this point, though the interested reader is invited to go to [probabilistic methods section] for much more on this perspective.

We note that in both cases of our argument, the lower bound for $R(n)$ we gave was non-constructive, i.e. we did not explicitly construct a coloring of K_N which avoids monochromatic K_n's, we only showed that such a coloring must exist. It is a major open problem to find a constructive argument which gives anywhere close to these bounds here.

Open Problem 3.8. *For some $c > 1$, find "explicit" red-blue edge colorings of K_{c^n} which avoid monochromatic K_n's.*

Observe that our proof not only shows that constructions should exist for $c = \sqrt{2}$, but in fact a more careful inspection shows that for any $c < \sqrt{2}$ that *almost every* coloring should work. Nevertheless, how to explicitly find such a coloring problem remains quite elusive.

The results we have mentioned in this sections are all classical, and the reader might wonder what is the current state of the art. For the lower bound, the only improvement over [result] is an argument due to Lovász using a slightly more involved probabilistic approach that gives a lower bound of [whatever], improving the bound of [result] by a multiplicative factor of [whatever].

For the upper bound, modest results showing bounds of the form $4^{n-o(n)}$ for an increasing series of $o(n)$ functions were obtained over the years until a recent major breakthrough by [authors in year] who proved that $R(n) \leq$???, and since then some further optimizations of their argument has yielded a bound of $R(n) \leq$???. At present this is all that is known for diagonal Ramsey numbers despite decades of hard work from an armada of talented mathematicians.

In addition to the diagonal Ramsey numbers $R(n)$, a lot of work has been put into studying the assymetric case $R(m, n)$. In particular, the study of these numbers when m is fixed and n tends towards infinity is referred to as "off-diagonal" Ramsey numbers. These problems are essentially equivalent to asking: how large can $\alpha(G)$ be if G is K_m-free and contains a given number of vertices? Indeed, more exposition, also comment on how $m = 3, 4$ are reasonably well understood due to complex probabilistic arguments.

3.2 More Colors and Arithmetic Ramsey Theory

There are a ton of variants for Ramsey numbers that one can consider. One of the immediate ones to consider is using more than just two colors. To this end, we define the *multi-color Ramsey number* $R_q(n)$ to be the smallest number N such that every q-coloring of the edges of K_N contains a monochromatic copy of K_n. Similar to [before] one can show that these numbers exist. In particular, we leave it as an exercise to prove the following bounds for the first non-trivial case of $n = 3$.

Theorem 3.9. *We have*

$$2^q < R_q(3) \leq 3 \cdot q!$$

Another direction is to consider coloring combinatorial objects other than graphs. One natural choice would be the integers $[N]$, from which we can ask if there exists a monochromatic subset satisfying some sort of arithmetic condition. One classical result due to Schur is as follows.

Theorem 3.10 (Schur)**.** *For all $q \geq 1$, there exists a finite number N_q such that any q-coloring of $[N]$ contains a monochromatic solution to the equation $x+y=z$, i.e. there exist three integers x, y, z with $x+y=z$ which are all assigned the same color.*

Proof. We will in fact prove that

$$N_q \leq R_q(3),$$

following a common theme in Ramsey theory of upper bounding one Ramsey problem by a function of another. To prove this, we will start with some coloring $\chi : [N] \to [q]$ and then use this to construct an auxiliary coloring $\chi' : E(K_N) \to [q]$ in such a way that monochromatic triangles under χ' correspond to monochromatic solutions to $x+y=z$ under χ. There are a couple of plausible ways one might try and define χ'. For example, given the edge $xy \in E(K_N)$ it is perhaps natural try coloring this edge to be the same color as either $\min(x, y)$ or $\max(x, y)$, but neither of these are really "compatible" with the goal of finding a solution to $x+y=z$.

With a bit more thought, one might come up with the (correct) idea of defining $\chi'(xy) = \chi(|x-y|)$. To see why this does what we want, assume that χ' has a monochromatic triangle on $u < v < w$. This implies that $\chi(v-u), \chi(w-v), \chi(w-u)$ all have the same color. Moreover, we have $(v-u)+(w-v)=(w-u)$, so taking $x = v-u,\ y = w-v$, and $z = w-u$ gives a monochromatic solution under χ. In total this implies that if $N \geq R_q(3)$ and χ is an arbitrary coloring then, because χ' must contain a monochromatic triangle since $N \geq R_q(3)$, χ contains a monochromatic solution to $x+y=z$. This proves $N_r \leq R_q(3)$, and in particular that this number is finite. $\square$

A lot more can be said about this area known as arithmetic Ramsey theory. Perhaps the most famous result in this direction is Van der Waerden's Theorem.

Theorem 3.11 (Van der Waerden's Theorem)**.** *For all k, q, there exists a finite number $N_{k,q}$ such that any r-coloring of $[N_{k,q}]$ contains a monochromatic k-term arithmetic progression. That is, there exist integers $a, d \geq 1$ such that $a, a+d, \ldots, a+(k-1)d$ are all given the same color.*

Proving this is not so easy, and the bounds for $N_{k,q}$ are horrendous even in the case of $q=2$. In fact, an even stronger statement than Van der Waerden's Theorem is known to be true.

Theorem 3.12 (Szemerédi's Theorem)**.** *Every subset of $[N]$ which does not contain a k-term arithmetic progression has size $o(N)$.*

To see this implication, observe that every r-coloring of $[N]$ contains a subset of size at least $r^{-1}N$ which, by Szermedi's Theorem, must contain a k-term arithmetic progression whenever N is sufficiently large. This is an example of a general phenomenon where Turán results (which bound how dense a structure can be before it contains a given substructure) often upper bound Ramsey results (which bound how large a structure can be with the property that it can be partitioned into r substructures avoiding a given substructure) simply because one of the partition elements in a Ramsey result must have relatively large density.

3.3 Ramsey Theory Without Colors

We will omit this for time unless requested by popular by demand. Broadly speaking it will be around the theme that Ramsey isn't just about saying that colored objects contain things. Some examples include monotone sequences and convex sets.

3.4 Exercises

1. Let's look at some small Ramsey numbers:

   (a) Prove that $R(3,4) = 9$ [2]

   (b) Prove that $R(4) \leq 18$ [1].

   (c) Prove that $R(4) = 18$ [3].

   (d) Determine[8] $R(5)$ [5].

2. Prove that every n-vertex graph has a clique or independent set on at least $\frac{1}{2}\log_2(n)$ vertices [1+].

3. Recall that a tournament is a digraph obtained by giving an orientation to each edge of a complete graph, and that a tournament is transitive if one can order its vertices $v_1, \ldots, v_n$ in such a way that $v_i \to v_j$ if and only if $i < j$. Prove that every tournament on n vertices contains a transitive tournament of size at least $\lfloor \log_2(n) \rfloor + 1$ [2-].

4. Here we sketch how to prove a lower bound for the first non-trivial offdiagonal Ramsey number $R(3,n)$.

   (a) Prove that to show $R(m,n) > N$ it suffices to construct a red-blue edge coloring of K_{2N} such that the number of red K_m's plus the number of blue K_n's is at most N [1+].

   (b) Prove that there exists $\varepsilon > 0$ such that $R(3,n) = \Omega(n^{1+\varepsilon})$. What's the best value of ε you can find using this method? (Hint: you will want to consider a random construction, but you'll want to color each edge red with some probability $p \ll \frac{1}{2}$ since there is assymetry in m and n) [2+].

* * *

5. Prove for all $n, q \geq 2$ that $R_q(n) \leq q^{qn}$ [2-].

[8] Currently the best known bounds are $43 \leq R(5) \leq 46$. The fact that this is still open should demonstrate how hard determining $R(n)$ exactly is. Indeed, Erdős once said something to the effect of: if aliens came to Earth and demanded we tell them what $R(5)$ was in the next 10 years or they would destroy us, then we should dedicate all our resources to this problem. If instead they ask for $R(6)$, then we should instead dedicate all our resources to fighting the aliens because we have no hope of doing what they ask.

6. Let us look at the multi-color Ramsey number $R_q(3)$.

   (a) Prove that $R_q(3) > 2^q$ [2-].

   (b) Prove that $R_q(3) \le 3 \cdot q!$, noting that this is best possible for $q = 2, 3$ [2].

   (c) Improve this upper bound to $R_q(3) \le \lfloor e \cdot q! \rfloor + 1$, which as far as we know is still the best known upper bound [3-].

7. For every graph F and integer q, define $R_q(F)$ to be the smallest integer N such that any q-edge coloring of K_N contains a monochromatic copy of F. Prove that if $\text{ex}(n, F) = O(n^{2-\alpha})$ for some $\alpha > 0$, then $R_q(F) = O(q^{1/\alpha})$ (Hint: concretely if you assume $\text{ex}(n, F) \le Cn^{2-\alpha}$ then you should be able to prove something like $R_q(F) \le (4Cq)^{1/\alpha}$, for example) [2-].

8. One of the most important results in general Ramsey theory is the Hales-Jewett Theorem which is a sort of "high-dimensional tic-tac-toe" theorem that goes as follows: Insert statement, exercise is to derive Van der Waerden from this.

* * *

9. We say that a graph G is K_n-Ramsey if any red-edge edge coloring of G contains a monochromatic copy of K_n, and we define the *size Ramsey number* $\hat{R}(n)$ to be the smallest number of edges in a graph which is K_n-Ramsey.

   (a) Observe that $R(n)$ can be defined to be the smallest number of *vertices* in a graph which is K_n-Ramsey, motivating this definition [1].

   (b) Prove that $\hat{R}(n) \le \binom{R(n)}{2}$ [1+].

   (c) Prove that $\hat{R}(n) = \binom{R(n)}{2}$; noting crucially that this equality holds despite us largely not understanding what $R(n)$ is (Hint: prove that any K_n-Ramsey graph must have chromatic number at least $R(n)$) [2+].

Part II

Structural Graph Theory

Insert flowery introduction.

4 Colorings

Recall that a *proper k-coloring* of a graph G is a map $c : V(G) \to [k]$ such that $c(u) \neq c(v)$ whenever $uv \in E(G)$. That is, we color each vertex using one of k colors such that no edge is monochromatic. We say that G is *k-colorable* if there exists a proper k-coloring of G, and we define the *chromatic number* $\chi(G)$ to be the smallest k such that G is k-colorable.

Colorings arise in various applied and theoretical contexts, and many problems in graph theory center around determining $\chi(G)$ for various graphs G. However, it is well known that determining whether $\chi(G) = k$ is an NP-hard problem for all $k \geq 3$, meaning one can not hope to find some "simple" way of determining if a graph has a given chromatic number. As such, the best one can realistically hope for in general is to establish reasonable bounds on $\chi(G)$ based on easy to compute parameters of G. We discuss two of the most fundamental bounds in the following sections.

4.1 Upper Bounds

Here and throughout this chapter we let $\Delta(G)$ denote the maximum degree of G, and whenever G is clear from context we will denote this quantity simply by Δ.

Theorem 4.1. *If G is a graph then*

$$\omega(G) \leq \chi(G) \leq \Delta(G) + 1.$$

Proof. We define a "greedy" coloring $c : V(G) \to [\Delta + 1]$ as follows. Let $v_1, \ldots, v_n$ be an arbitrary ordering of the vertices of G. Iteratively given that we have defined $c(v_1), \ldots, c(v_{i-1})$ we choose $c(v_i)$ to be any element in $[\Delta + 1] \setminus \{c(v_j) : v_j \in N(v_i),\ j < i\}$; note that such an element must exist since $|N(v_i)| < \Delta + 1$.

We claim that c is a proper $(\Delta + 1)$-coloring. Indeed, if $v_i v_j \in E(G)$ with, say, $i > j$ then we chose $c(v_i)$ to be disjoint from $c(v_j)$. Thus c is a proper $(\Delta + 1)$-coloring, proving the result. $\square$

Theorem 4.1 is important in the field of coloring because it and its proof serves as the starting point for a number of other foundational results, several of which we discuss now.

Perhaps the immediate question to ask upon seeing Theorem 4.1 is if this bound is tight. And indeed, one quickly sees that it is for $K_{\Delta+1}$ for all $\Delta \geq 1$, and for $\Delta = 2$ it is tight if and only if G contains an odd cycle. This turns out to exactly describe the cases of equality for Theorem 4.1.

Theorem 4.2 (Brooks's Theorem)**.** *If G is a connected graph of maximum degree Δ and if G is not an odd cycle or $K_{\Delta+1}$ then $\chi(G) \leq \Delta$.*

Sketch of Proof. Essentially one can show that if G is as in the hypothesis, then there exists an ordering $v_1, \ldots, v_n$ of $V(G)$ such that (1) $v_1, v_2 \in N(v_n)$, (2) $v_1 \not\sim v_2$, and (3) $|\{v_j \in N(v_i) : j < i\}| < \Delta$ for all $i < n$. We now consider a greedy coloring $c : V(G) \to [\Delta]$ as we did before except (crucially) we set $c(v_1) = c(v_2)$ which will not create an improper coloring by (2). By

(3), every vertex $v_i < n$ will have at least 1 choice when it is time to be colored, and by (1) the set $\{c(v_j) : v_j \in N(v_n)\}$ has at most $\Delta - 1$ used colors since $c(v_1) = c(v_2)$, meaning that we can also color v_n successfully. This gives a proper Δ-coloring of G as desired. □

While the maximum degree of a graph is a nice, clean parameter, it can often be entirely unrelated to $\chi(G)$ with perhaps the most egregious example of this being the star $K_{1,\Delta}$ which has chromatic number 2. Given this, one can ask if its possible to strengthen the bound of Theorem 4.1 by using some sort of "refinement" of the maximum degree Δ which, in particular, gives more reasonable bounds for stars. To this end and with Mantra 9 as motivation, we might ask ourselves what the best possible bound we could prove using the same argument as in Theorem 4.1, giving rise to the following parameter.

Definition 8. We define the *degeneracy* of a graph G to be the smallest integer d such that there exists an ordering $v_1, \ldots, v_n$ of $V(G)$ such that $|\{v_j \in N(v_i) : j < i\}| \leq d$ and we denote the degeneracy of G by $d(G)$.

With this definition the exact same proof of Theorem 4.1 gives the following.

Theorem 4.3. *If G is a graph, then*

$$\chi(G) \leq d(G) + 1.$$

Note that we always have $d(G) \leq \Delta$ via considering an arbitrary ordering of $V(G)$, so Theorem 4.1 is always at least as strong as Theorem 4.1. Moreover, it is an exercise to show that $d(G) = 1$ whenever G is a forest with at least 1 edge, meaning Theorem 4.3 is tight for all such graphs.

As an aside, the reader might feel that our definition of degeneracy is rather ad-hoc and specific only to the very particular proof we were trying to generalize. However, it turns out that degeneracy plays an important role in other areas such as Turán problems and that it has other (perhaps more natural) equivalent formulations. We touch on some of these connections in the exercises.

The last extension of Theorem 4.1 that we touch on asks if we can not only find some proper $(\Delta + 1)$-coloring but one which has some additional "nice" properties. This is perhaps natural to consider given that a closer look at our proof of Theorem 4.3 reveals that there is not just one proper $(\Delta + 1)$-coloring but in fact exponentially many, so we can perhaps be a bit more greedy with the sort of coloring we get at the end. There are various "nice" properties one could consider for colorings; the one we focus on will be the following.

Definition 9. We say that a proper k-coloring $c : V(G) \to [k]$ is *equitable* if $|c^{-1}(i)| \in \{\lfloor v(G)/k \rfloor, \lceil v(G)/k \rceil\}$ for all i. That is, each color is used as equal a number of times as possible.

Equitable colorings are a lot harder to come by compared to usual colorings. Indeed, the star $K_{1,\Delta}$ has exponentially many proper 3-colorings but none of them are equitable if $\Delta \geq 5$. However, it turns out that equitable colorings always exist at the threshold of $\Delta + 1$.

Theorem 4.4 (Hajnal-Szemerédi Version II)**.** *If G is a graph with maximum degree Δ then there exists an equitable proper $(\Delta + 1)$-coloring.*

This result turns out to be equivalent to our previous statement Theorem 2.10 of the Hajnal-Szemerédi Theorem, which is perhaps surprising at first glance but which is not too hard to prove; we leave this as an exercise. As before, we refrain from proving this result.

4.2 Lower Bounds and Perfect Graphs

For our lower bounds we recall that $\alpha(G)$ denotes the largest size of an independent set of G and that $\omega(G)$ denotes the largest size of a clique of G.

Theorem 4.5. *For every graph G, we have*

$$\chi(G) \geq \omega(G),$$

and

$$\chi(G) \geq \frac{v(G)}{\alpha(G)}.$$

Proof. The first bound follows simply because the vertices making up the clique of size $\omega(G)$ of G must all be given colors that are distinct from each other. For the second bound, we observe that in any proper coloring $c : V(G) \to [t]$ that $c^{-i}(i)$ is an independent set of G for all i (otherwise c would have two adjacent vertices mapped to the same color i). In particular, one of these independents has size at least $v(G)/t$ by the pigeonhole principle, and taking $t = \chi(G)$ implies that $\alpha(G)$ is at least $v(G)/\alpha(G)$ as desired. $\square$

Both of these bounds can easily seen to be tight for $G = K_n$. However, characterizing all cases of equality analogous to Brooks's Theorem seems difficult to do here. Indeed, $\chi(G) = v(G)/\alpha(G)$ holds if and only if $V(G)$ has a partition into maximum independent sets and offhand there does not seem to be a simple way to characterize this property. The case of $\chi(G) = \omega(G)$ is even more complex, as for any graph G' we can form a graph $G = G' \sqcup K_n$ with $n = v(G')$ and this trivially satisfies $\chi(G) = \omega(G)$ despite the structure of G being entirely arbitrary on half of its vertices. To avoid having to take into account silly constructions like these, we will want to shift to studying a certain class of graph families which are ubiquitous in structural graph theory.

Definition 10. We say that a family of graphs $\mathcal{G}$ is *hereditary* if it is closed under deleting vertices, that is, if for every $G \in \mathcal{G}$ we have $G - v \in \mathcal{G}$ for every vertex $v \in V(G)$. Equivalently, $\mathcal{G}$ is hereditary if for every graph $G \in \mathcal{G}$ all of the induced subgraphs of G are also in $\mathcal{G}$.

Many natural families of graphs are hereditary, such as those avoiding some graph F as either an induced or non-induced subgraph. Returning to our previous problem, we will now aim to characterize not the full family of graphs $\mathcal{G}$ with $\chi(G) = \omega(G)$ for all $G \in \mathcal{G}$ but simply the largest hereditary family of graphs $\mathcal{G}$ with this property. Equivalently, we aim to study the following type of graphs.

Definition 11. We say that a graph G is *perfect* if $\chi(G') = \omega(G')$ for every induced subgraph G' of G.

Again to be clear, the family of all perfect graphs is a hereditary family and it is the largest one satisfying $\chi(G) = \omega(G)$ for every graph in the family. Perfect graphs have a long history of study with this ultimately culminating in a full characterization of their structure.

Theorem 4.6 (Strong Perfect Graph Theorem)**.** *A graph G is perfect if and only if both G and its complement $\overline{G}$ do not contain an induced odd cycle of length at least 5.*

The fact that graphs must satisfy this proper to be perfect is an exercise. The converse is tremendously difficult to prove and was originally done so by Chudnovsky, Robertson, Seymour, and Thomas in 2006.

Possibly do a proof of weak perfect graph theorem.

4.3 Coloring Variants

Here we look at some variants of the notion of proper colorings.

4.3.1 List Colorings

It is very common in coloring arguments to construct some proper k-coloring $c : V(G) \to [k]$ by inductively defining $c(v)$ for some vertex v and then constructing a coloring of $G - v$. However, when we do this we are no longer exactly looking for a proper k-coloring of $G - v$ but rather a coloring where each $u \notin N(v)$ is allowed to be any color in $[k]$ while $u \in N(v)$ are required to be colored from the set $[k] \setminus \{c(v)\}$, and because of this we can't directly apply any inductive statement that holds for proper k-colorings. The solution to this problem is to consider a more general notion of coloring which is preserved by us iteratively coloring a vertex of our graph. Specifically, we do this by assigning each vertex a list of "allowed colors" $L(v)$ which we can think of as being the subset of $[k]$ obtained after removing any of the colors from vertices we've already deleted from G in some sort of inductive step. More precisely, we have the following.

Definition 12. Given a graph G, a *list assignment* is a function L which assigns to each $v \in V(G)$ a set $L(v)$. A *proper L-coloring* is a map c from $V(G)$ which satisfies $c(v) \in L(v)$ for all $v \in V(G)$ and which has $c(u) \neq c(v)$ for all u, v with $uv \in E(G)$. We say that G is *k-choosable* if there exists a proper L-coloring for G for all L with $|L(v)| \leq k$ and we define the *list chromatic number* $\chi_\ell(G)$ to be the smallest k such that G is k-choosable.

As an example, observe that G has a proper k-coloring if and only if it has a proper L-coloring with $L(v) = [k]$ for all v. As such, G being k-choosable implise that it is k-colorable and hence $\chi(G) \leq \chi_\ell(G)$ for every graph G. As such, the following is a direct strengthening of the results of the previous subsection.

Theorem 4.7. *If G is a graph of maximum degree Δ then*

$$\chi_\ell(G) \leq d(G) + 1 \leq \Delta + 1.$$

The proof of this is essentially identical to our previous arguments and we leave the details as an exercise to the reader.

While Theorem 4.7 is certainly at least as strong as our results upper bounding $\chi(G)$, it is not clear if this is a strict strengthening. That is, it is not clear whether there exists any graph with $\chi(G) \neq \chi_\ell(G)$. And indeed, intuitively it doesn't like this should be the case. That is, finding a proper L-coloring seems hardest to do when the lists $L(v)$ overlap as much as possible since otherwise it seems easier for us to avoid creating monochromatic edges. As such, it naively seems like the worst-case scenario for L is if $L(v) = [k]$ for all v which exactly recovers the notion of a proper k-coloring.

Perhaps surprisingly (or unsurprising given we've dedicated a whole subsubsection to this topic), there do in fact exist L which are strictly harder to properly color compared to the identically $[k]$ assignment, implying that $\chi(G) < \chi_\ell(G)$ for such graphs. Genuinely surprisingly, this holds even for bipartite graphs where $\chi(G)$ and $\chi_\ell(G)$ can be made arbitrarily far apart from each other.

Theorem 4.8. *For every integer $t \geq 2$, there exists a graph G with $\chi(G) = 2$ and $\chi_\ell(G) \geq t$.*

Sketch of Proof. We only prove this for $t = 3$ with the generalization of this argument being left as an exercise to the reader. For this, take $G = K_{3,3}$ say with bipartition $U = \{u_1, u_2, u_3\}$ and $V = \{v_1, v_2, v_3\}$. Define L by having $L(u_i) = L(v_i) = \{1, 2, 3\} \setminus \{i\}$. Observe now that if there exists a proper L-coloring c then $\{c(u_1), c(u_2), c(u_3)\}$ contains at least 2 colors since if this only contained one color i then this would contradict $c(u_i) \in L(u_i) = \{1, 2, 3\} \setminus \{i\}$. But this set containing at least two colors implies that $L(v_i) \subseteq \{c(u_1), c(u_2), c(u_3)\}$ for some i, namely the one whose color set $L(v_i)$ equals these two colors. This means that for any choice of $c(v_i) \in L(v_i)$ that there will exist some $u_j \in N(v_i)$ with $c(u_j) = c(v_i)$, contradicting this color being proper. We conclude that no proper L-coloring can exist for this choice of L, implying that $\chi_\ell(K_{3,3}) > 2$. □

Corollary 4.9. *There does not exist a function $f : \mathbb{N} \to \mathbb{N}$ such that $\chi_\ell(G) \leq f(\chi(G))$ for all graphs G. That is, the list-chromatic number can not be bounded by some function of the chromatic number.*

As a final remark, we note that in very recent years an even greater generalization of list coloring has appeared in the literature known alternatively as *correspondence coloring* or *DP-coloring* which in particular originated in the context of inductive proofs similar to our motivation for studying list colorings. Maybe say more about this at some point.

4.3.2 Edge Colorings

All of the colorings we have considered up to this point involve colorings of the vertices of G. What if we were to consider colorings of its edges instead? While a natural idea, it is not immediately clear what a "proper" edge coloring should be. Motivated by the idea that a vertex coloring is proper if no two vertices which share an edge in common are given the same color, we might consider edge colorings to be proper if no two edges which share a vertex in common are given the same color.

Definition 13. Given a graph G, we say that a function $c' : E(G) \to [k]$ is a *proper k-edge coloring* if $c'(e) \neq c'(f)$ for any distinct edges e, f with $e \cap f \neq \emptyset$. We define the *chromatic index* $\chi'(G)$ to be the smallest integer k such that G has a proper k-edge coloring.

Proper edge colorings of a graph G are in fact equivalent to proper vertex colorings of an appropriate auxiliary graph of G.

Lemma 4.10. *Given a graph G, define the line graph $L(G)$ to be the graph with vertex set $E(G)$ where two distinct edges e, f are adjacent to each other in $L(G)$ if $e \cap f \neq \emptyset$. A function $c' : E(G) \to [k]$ is a proper k-edge coloring of G if and only if it is a proper k-coloring of $L(G)$.*

We can use this connection to proper colorings to immediately conclude some very strong bounds on $\chi'(G)$.

Proposition 4.11. *If G is a graph with maximum degree Δ, then*

$$\Delta \leq \chi'(G) \leq 2\Delta - 1.$$

Proof. Indeed, observe that $\omega(L(G)) \geq \Delta$ as the Δ edges incident to a vertex of maximum degree in G form a clique in the line graph $L(G)$. On the other hand, the maximum degree of $L(G)$ is at most $2\Delta - 2$ as every edge uv in G is incident to at most $2\Delta - 2$ edges other than uv itself (since each of u, v are incident to at most $\Delta - 1$ other edges respectively). The bounds now follow immediately from Theorem 4.1. $\square$

One can use Brooks's Theorem to improve the upper bound of this proposition by 1 for $\Delta \geq 3$, but we choose not to do so here since a substantially stronger bound holds.

Theorem 4.12 (Vizing's Theorem)**.** *If G is a graph with maximum degree Δ then $\chi'(G) \in \{\Delta, \Delta + 1\}$.*

We omit the proof due to Guantao literally teaching a full course on edge colorings right now.

Despite Vizing's Theorem determining χ' up to an additive error of 1 for every graph G there is still a lot that can be said about edge colorings especially in the context of multigraphs, though we will not go into this further here.

4.3.3 Fractional Colorings

We won't cover due to time constraints but roughly the idea is to take a fractional relaxation of the coloring integer programming problem via allowing vertices to be 2/3rd colored red and 1/3rd colored blue.

4.4 Clique Numbers and Chromatic Numbers

A major theme of structural graph theory is to determine when a given parameter of a graph G can be bounded by a function of another parameter. For example, we saw that $\chi_\ell(G)$ can not be upper bounded by a function of its natural lower bound $\chi(G)$ while $\chi'(G)$ can be very strongly upper bounded by its natural lower bound $\Delta(G)$. For $\chi(G)$, the natural question to ask in view of Theorem 4.5 is whether $\chi(G)$ is upper bounded by a function of its clique number $\omega(G)$. As a first step, we need to figure out if an analog of Theorem 4.8 holds in our setting.

Question 4.13. *Is it true that for every t there exists a graph G with $\omega(G) = 2$ but $\chi(G) \geq t$?*

That is, do there exist triangle-free graphs with arbitraily large chromatic numbers? The answer to this question is immediately yes for $t = 3$ by considering odd cycles. One can also verify it for $t = 4$, though it likely will take you either a lot of trial and error or a computer (as the smallest such example is on 11 vertices), which are approaches which will not generalize to, say, $t = 1000$. The difficulty in finding these constructions should suggest that either this is false for large t or that we need a more systematic scheme for forming our constructions. And indeed, we will in fact show that this question has a positive answer by coming up with a systematic way for constructing examples.

The motivation for our approach is as follows. Say we have some triangle-free graph G with chromatic number at least t, we want to build from this a new graph $M(G)$ which is triangle-free and which has chromatic number at least $t+1$. The simplest way to force chromatic number at least $t+1$ is to add a new vertex w to G which is adjacent to all of $V(G)$ since the new vertex is forced to be given a coloring distinct from the t which we know must be used for G, but this approach completely fails to maintain that our graph is triangle-free. To get around this, for each $u_i \in V(G)$ we will create a new "duplicate" vertex v_i in such a way that we essentially force the color of v_i to be the same as the color of u_i and such that these duplicate vertices v_i form an independent set. If we can achieve this, then by adding a new vertex w adjacent to all of the duplicate vertices will achieve our desired goal. After pondering on this idea for a bit one might be led to the following operation.

Definition 14. Given a graph G with vertices $u_1, \ldots, u_n$, its *Mycielskian* $M(G)$ is a graph with vertex set $u_1, \ldots, u_n, v_1, \ldots, v_n, w$ such that:

- $u_iu_j \in E(M(G))$ and $u_iv_j \in E(M(G))$ if and only if $u_iu_j \in E(G)$,
- $v_iw \in E(M(G))$ for all i, and
- $u_iw \notin E(M(G))$ and $v_iv_j \notin E(M(G))$ for all i, j.

and such that $u_iw \notin E(M(G))$ and

Insert picture of $G = K_2$ and also maybe $G = C_5$..

That is, $M(G)$ is formed by taking G, duplicating each vertex so that v_i has the same set of neighbors as u_i in G, and then adding a new vertex w adjacent to all the duplicated vertices. Crucially, this operation does precisely what we want it to do.

Proposition 4.14. *For every graph G, $\chi(M(G)) = \chi(G) + 1$ and $M(G)$ is triangle-free whenever G is triangle-free.*

Proof. For triangle-freeness, we observe that no triangle in $M(G)$ can involve two v_i vertices since such vertices are never adjacent, and as such no triangle can involve w whose only neighbors are v_i vertices. As such, if there is a triangle it must either be of the form u_i, u_j, u_k or v_i, u_j, u_k, but such vertices form a triangle in $M(G)$ if and only if u_i, u_j, u_k form a triangle in G, proving this half of the result.

For ease of notation let $t = \chi(G)$. To prove $\chi(M(G)) \leq t+1$ we construct an explicit proper $(t+1)$-coloring for $M(G)$ as follows. Start with an arbitrary proper t-coloring c' of G. Now define $c : V(M(G)) \to [t+1]$ by having $c(u_i) = c(v_i) = c'(u_i)$ and $c(w) = t+1$. That is, we duplicate the coloring of c' on both the u and v vertices and then give w a completely new color. Any edge involving w will be monochromatic because w is the only vertex with color $c(w)$. One can also check that if some edge u_iu_j or v_iv_j were monochromatic under c then the edge u_iu_j would be monochromatic under c' which we assumed not to be the case. This shows c is a proper coloring, proving the bound.

We now prove that $\chi(M(G)) \geq t+1$, and for this we assume for contradiction that there exists some proper t-coloring c of $M(G)$.

Claim 4.15. *For every color $s \in [t]$, there exists some u_i with $\{c(u_j) : u_j \in N_G(u_i)\} = [t] \setminus \{s\}$.*

Proof. Assume this was false for some s, we aim to use this to contradict that G has chromatic number t. To this end, define a coloring $c' : V(G) \to [t] \setminus \{s\}$ by having $c'(u_i) = c(u_i)$ whenever $c(u_i) \neq s$ and otherwise take $c'(u_i)$ to be an arbitrary color in $[t] \setminus (\{s\} \cup \{c(u_j) : u_j \in N_G(u_i)\})$, noting that such a color exists by hypothesis. We claim that this is a proper coloring. Indeed, the only way an edge u_iu_j can be monochromatic under c' is if, say, $c(u_i) = s$, but in this case we must have $c(u_j) \neq s$ since c is proper coloring and hence $c'(u_j) = c(u_j) \neq c'(u_i)$ by construction. We have thus shown that G can be properly colored using only $t-1$ colors, contradicting $\chi(G) = t$. □

With this claim we see that $\{c(v_1), \ldots, c(v_n)\} = [t]$ since for each u_i as in the claim we must have $c(v_i) = s$. But this means $c(w)$ will equal the color of one of its neighbors v_i, a contradiction to c being a proper coloring. □

Corollary 4.16. *For all $t \geq 2$ there exists a triangle-free graph with chromatic number t.*

Proof. Take $G_2 = K_2$ and iteratively define $G_{i+1} = M(G_i)$. The proposition immediately implies that G_t satisfies the conditions of the corollary. □

A natural followup now is to ask to what extent we can strengthen this result. For example, what if our graph is both C_3-free and C_5-free (the two smallest certificates for whether a graph has chromatic number 2 or not), can we find graphs of arbitrarily large chromatic number in this case? Note here that the Mycielskian $M(G)$ will be entirely ineffective for this sort of problem since any edge in G creates a C_5 in $M(G)$. It is natural then to go back to our motivation for $M(G)$ and see if one can modify it to get rid of C_5's as well, but we do not know of any way to make this work.

Ultimately there does exist an explicit construction of graphs with no C_3, C_4, or C_5 which have arbitrarily large chromatic numbers due to Tutte. However, these graphs are tremendously large and as far as we know these constructions do not generalize to the next natural followup question of asking if there exist graphs with large chromatic number which avoid all of C_3, C_5, and C_7. Ultimately, this problem does indeed have a positive answer in a very strong sense. To this end, we recall that the *girth* of a graph is the length of its shortest cycle.

Theorem 4.17 (Erdős)**.** *For all integers $\ell, t \geq 2$ there exists a graph G with girth at least ℓ and $\chi(G) \geq t$.*

This result says in a very strong sense that $\chi(G)$ is a "global" parameter of G, in the sense that it implies there exist graphs which G locally look like a tree (in the sense that G restricted to the vertices within distance $g/2$ of a given vertex is a tree) but nevertheless needs an arbitrarily large number of colors to actually color the whole graph.

There is no known family of "elementary" graphs[9] which satisfies Theorem 4.17. However, similar to our proof of Theorem 3.7 showing $R(n)$ is large we will be able to prove with an appropriate (though somewhat more involved) random construction. For this, we recall Markov's inequality which says that if X is a non-negative random variable then $\Pr[X \geq t] \leq \mathbb{E}[X]/t$ for all real t.

Proof. The random construction we consider will be based off perhaps the most important object in probabilistic combinatorics, namely the Erdős-Renyi random graph model. To this end, for an integer $n \geq 1$ and a real number $0 \leq p \leq 1$ we let $G_{n,p}$ denote the random n-vertex graph obtained by including each edge independently and with probability p. Thus $G_{n,1} = K_n$ with probability 1, $G_{n,0}$ is the empty graph with probability 1, and $G_{n,1/2}$ is equally likely to be any n-vertex graph. The naive idea we want to try is to pick some values for n and p such that with high-probability $G_{n,p}$ simultaneously has few (or even 0) cycles of length at most ℓ and has large chromatic number. Let us address each of these obstacles in turn.

First of all, let X_i denote the number of cycles of length i in G and let $X_{<\ell} = \sum_{i=3}^{\ell-1} X_i$. Observe that $\mathbb{E}[X_i] \leq p^i n^i$ as the total number of cycles of length i in K_n is at most n^i and the probability that any given cycle C survives into $G_{n,p}$ is exactly p^i (i.e. this is the probability that $G_{n,p}$ independently keeps all i edges of C). By linearity of expectation we find that

$$\mathbb{E}[X_{<\ell}] \leq \sum_{i=3}^{\ell-1} p^i n^i \leq (\ell-1)\max\{pn, (pn)^{\ell-1}. \tag{2}$$

We now turn to studying $\chi(G_{n,p})$, and a priori it is not so clear how to approach this. The key insight is that we only care about proving lower bounds for this chromatic number, so it suffices to bound some general lower bound for χ which might be simpler to analyze. In particular, by Theorem 4.5 it suffices to show that $n/\alpha(G_{n,p})$ is large, i.e. that $\alpha(G_{n,p})$ is small, which is much simpler to do. Indeed, if we let Y_a denote the number of independent sets of $G_{n,p}$ of size a then by linearity of expectation and the basic inequality $1 - x \leq e^{-x}$, we find

$$\mathbb{E}[Y_a] = (1-p)^{\binom{a}{2}} \cdot \binom{n}{a} \leq e^{-p\binom{a}{2}} \cdot n^a = (ne^{-p(a-1)/2})^a \leq (ne^{-pa/2})^a. \tag{3}$$

[9] There do exist explicit constructions due to Lubotzky, Phillips, and Sarnak, but these are highly complicated and rely on quite a bit of algebra and number theory.

Heuristically, what this tells us is that if $e^{-pa/2} \ll n^{-1}$, i.e. if $p \gg \frac{\log(n)}{a}$, then $G_{n,p}$ with high probability will not contain any independent sets of size at least a, which if true would imply that $\chi(G_{n,p}) \geq n/a$. This gives a good lower bound if $a \ll n$, and hence heuristically we need $p \gg \log(n)/n$ in order for us to conclude that $G_{n,p}$ has large chromatic number. Unfortunately though for this range of p, (2) suggests that the number of short cycles in $G_{n,p}$ could be as large as $(\log n)^\ell$. In total this suggests (the true fact that) $G_{n,p}$ does not simultaneously have high girth and high chromatic number for any choice of p.

The saving grace to this approach is the observation that although $G_{n,p}$ does not have 0 short cycles for $p \gg \log n/n$, it does have *few* of them. In particular, if we take $G_{n,p}$ and delete a vertex from each of its short cycles then this graph will by definition have large girth and also be very close to $G_{n,p}$ if $G_{n,p}$ has few short cycles.

With all this motivation in mind, let $p = C\log n/n$ with n, C sufficiently large integers so that the following inequalities hold. Let A_1 be the event that $X_{<\ell} \leq n/2$. By Markov's inequality and (2) we have

$$\Pr[A_1] = 1 - \Pr[X_{<\ell} > n/2] \geq 1 - \frac{(\ell-1)(pn)^{\ell-1}}{n/2} \geq 1 - \frac{(\ell-1)C^{\ell-1}(\log n)^{\ell-1}}{n/2} > \frac{1}{2},$$

with this last inequality holding for n sufficiently large in terms of C and ℓ. Similarly let A_2 denote the event that $\alpha(G_{n,p}) < n/2t$. By Markov's inequality and (3) we have

$$\Pr[A_2] = 1 - \Pr[Y_{n/2t} \geq 1] \geq 1 - (ne^{-pn/4t})^{n/2t} = 1 - (n \cdot n^{-C/4t})^{n/2t} > \frac{1}{2},$$

with the last inequality holding for $C > 4t$ and n sufficiently large. From this, we conclude that $\Pr[A_1 \cap A_2] > 0$, i.e. with positive probability both A_1 and A_2 occur, i.e. there exists an n-vertex graph G such that it has at most $n/2$ cycles of length less than ℓ and $\alpha(G) < n/2t$. Define G' by taking G and deleting 1 vertex from each cycle of length less than ℓ. By assumption we have $v(G') \geq n/2$ and $\alpha(G') \leq \alpha(G) \leq n/2t$, and as such

$$\chi(G') \geq \frac{v(G')}{\alpha(G')} \geq \frac{n/2}{n/2t} = t,$$

proving the result. $\square$

At this point we have more than proved that one can not upper bound $\chi(G)$ by a function of $\omega(G)$ for arbitrary graphs G, but what about if we turn our attention away from all graphs and restrict to some nice family of graphs instead?

Definition 15. We say that a family of graphs $\mathcal{G}$ is χ-*bounded* if there exists a function $f : \mathbb{N} \to \mathbb{N}$ such that $\chi(G) \leq f(\omega(G))$ for all $G \in \mathcal{G}$.

For example, Theorem 4.17 says that the family of all graphs is not χ-bounded. On the other hand, the family of perfect graphs by definition are χ-bounded with $f(n) = n$, and Vizing's Theorem implies that the family of line graphs is χ-bounded with respect to $f(n) = n + 1$. There are many open questions regarding which families of graphs are χ-bounded as well as determining optimal values for the function f with the biggest open question being the following.

Conjecture 4.18 (Gyárfás-Sumner)**.** *For every tree T, the family of graphs $\mathcal{G}_T$ which do not contain an induced copy of T is χ-bounded.*

This conjecture was originally made by Gyárfás in 1975 (and independently by Sumner later) and despite receiving a lot of attention, the only trees we know of for which this problem is solved is when T is a star, path, or has radius 2.

4.5 Exercises

1. We begin with some warmups.

   (a) Recall that a map $\phi : V(G) \to V(H)$ is a homomomorphism if $\{\phi(u), \phi(v)\} \in E(H)$ whenever $\{u, v\} \in E(G)$. Prove that a graph G has a proper t-coloring if and only if there exists a homomorphism $\phi : V(G) \to K_t$ [1].

   (b) Prove that if $c : V(G) \to [t]$ is a proper coloring then $c^{-1}(i)$ is an independent set of G. Prove that if $c' : E(G) \to [t]$ is a proper edge coloring then $(c')^{-1}(i)$ is a matching of G [1].

* * *

2. Prove that the number of proper t-colorings of a graph G is at least $\prod_{v\in V(G)}(t - \deg(v))$. In particular, every n-vertex graph G with maximum degree Δ has at least 2^n proper $(\Delta + 2)$-colorings [1+].

3. Prove that there exists some $C > 1$ such that every n-vertex graph G with maximum degree Δ has at least C^n proper $(\Delta+1)$-colorings. That is, G does not just have 1 proper $(\Delta + 1)$-coloring but at least exponentially such colorings (Hint: the proof we have in mind yields that there are at least $2^{\frac{\Delta}{\Delta+1}n}$ such colorings, in particular giving the result with $C = 2^{2/3}$) [2+].

4. Prove that one can equivalently define the degeneracy $d(G)$ of a graph G to be the smallest integer d such that every subgraph $G' \subseteq G$ contains a vertex of degree at most d [1+].

5. Prove that $e(G) \leq d(G)v(G)$ for all graphs G. In particular, graphs with bounded degeneracy have at most a linear number of edges [1+].

6. Let us consider the degeneracy of various types of graphs.

   (a) Prove that a graph G has $d(G) = 1$ if and only if G is a forest with at least one edge [1+].

   (b) Prove that if G is a graph with maximum degree Δ then $d(G) = \Delta$ if and only if G is regular. In particular, note that this gives an easy proof of Brooks's Theorem for graphs G which are not regular [1+].

(c) Prove that if G is a planar graph then $d(G) \le 5$. In particular, note that this implies $\chi(G) \le 6$ for planar graphs (Hint: you may assume without proof the fact that planar graphs have $e(G) \le 3v(G) - 6$ provided $v(G) \ge 3$) [1+].

7. Here we briefly showcase how degeneracy appears in other graph theoretic contexts.

   (a) Prove that for all d there exists a graph F with $d(F) = d$ such that $\mathrm{ex}(n, F) = \Theta(n^{2-1/d})$ (Hint: you may assume without proof anything I claimed in the chapter on Turán Problems) [2-].

   (b) Prove that if F is a graph with $d(F) = d$, then $\mathrm{ex}(n, F) = O(n^{2-\frac{1}{4d}})$; a major open problem of Erdős conjectures that in fact $\mathrm{ex}(n, F) = O(n^{2-\frac{1}{d}})$ should hold, which is best possible by the previous part [3+].

   (c) Recall that $R_2(F)$ denotes the smallest number N such that any red-blue edge coloring of K_N contains a monochromatic copy of F. Prove that if $d(F) = d$ then $R_2(F) = O_d(v(F))$ [4-].

8. Prove that our two stated versions of the Hajnal-Szemerédi Theorem are equivalent to each other [2-].

* * *

9. Prove that if a graph G is perfect, then G and $\overline{G}$ do not contain an induced odd cycle of length at least 5 [2].

10. Determine which of the following families of graphs are hereditary: all graphs, regular graphs, planar graphs, trees, forests [1].

11. Prove that for every hereditary family $\mathcal{G}$ that there exists a family of graphs $\mathcal{F}$ such that $G \in \mathcal{G}$ if and only if G does not contain any graph of $\mathcal{F}$ as an induced subgraph. Prove that there exists a hereditary family $\mathcal{G}$ such that the family $\mathcal{F}$ can not be taken to be finite [2-].

* * *

12. Formally prove that $\chi_\ell(G) \le d(G) + 1$ for every graph G [1].

13. We consider the list chromatic number of complete bipartite graphs.

   (a) Complete our proof of Theorem 4.8 by showing that for all t there exists some n_t such that $\chi_\ell(K_{n_t,n_t}) \ge t$ [2].

   (b) Give an alternative proof by showing that $\chi_\ell(K_{t-1,(t-1)^{t-1}}) \ge t$ [2].

14. Use Brooks's Theorem to prove that $\chi'(G) \leq 2\Delta - 2$ for every graph G with maximum degree $\Delta \geq 3$ [1+].

$* * *$

15. Prove that every K_3-free graph on at most 10 vertices has chromatic number at most 3, meaning that $M(C_5)$ is the smallest triangle-free graph with chromatic number 4. Your proof should be human readable and checkable, i.e. it can not be of the form "I generated every graph on at most 10 vertices on my computer and verified that this is true" [3-].

16. Given a graph G and an integer $k \geq 1$, we define the Zykov graph $Z(G, k)$ by taking k disjoint copies $G_1, \ldots, G_k$ of G, and then for each of the $v(G)^k$ sequences $\vec{x} = (x_1, \ldots, x_k) \in V(G_1) \times \cdots \times V(G_k)$ we add a new vertex $v_{\vec{x}}$ whose neighborhood equals $\{x_1, \ldots, x_k\}$.

    Prove that $Z(G, k)$ is triangle-free whenever G is, and that $\chi(Z(G, k)) = \chi(G)+1$ provided $k \geq \chi(G)$. As such, these graphs give another explicit family of triangle-free graphs with arbitrarily large chromatic numbers [2].

17. Prove Markov's inequality whenever X is a non-negative discrete random variable [1+].

18. In this exercise we will partially motivate the exact statement of the Gyárfás-Sumner conjecture. To this end, for each graph F let $\mathcal{G}_F$ be the family of graphs which does not contain F as an induced subgraph.

    (a) Prove that if F contains a cycle then $\mathcal{G}_F$ is not χ-bounded [2-].

    (b) Prove that if F_1, F_2 are forests then $F_1 \sqcup F_2$ is χ-bounded if and only if F_1 and F_2 are both χ-bounded (Hint: inductively define $f(1)$, then $f(2)$, then $f(3)$, and so on) [2+].

    (c) Conclude that to determine which graphs F are such that $\mathcal{G}_F$ is χ-bounded it suffices to do so in the case when F is a tree [1].

5 Matchings and Factors

I have been told you learned most of this in 6420 so I'll just give a quick recap and prove Tutte unless I am told otherwise.

In Section 2 we saw some sufficient conditions for G to contain a Hamiltonian cycle, and it is natural to ask if there exists a nice necessary and sufficient condition for Hamiltonicity. This turns out to be essentially hopeless. Indeed, it is known that the computational problem of determining whether or not a given graph is Hamiltonian is NP-complete, which means that if a "simple" necessary and sufficient condition existed then a large number of seemingly intractable problems for computer science would all have efficient algorithms. Similarly determining whether a graph has a Hamiltonian path is also NP-complete. However, there does exist a nice characterization for when a graph has a perfect matching.

Definition 16. Given a graph G, a *matching* M is a subgraph of G such that every vertex has degree 1, i.e. M is the disjoint union of some number of edges. The *matching number* $\nu(G)$ is the maximum number of edges in a matching of G. A *perfect matching* is a matching which is incident to every vertex of G. Note that a perfect matching can only exist if $v(G)$ is even.

Again this is being rushed because I assume we've seen it already.

Matchings are particularly nice in bipartite graphs where we have the following two fundamental (and ultimately equivalent) theorems.

Theorem 5.1 (Kőnig's Theorem)**.** *Given a graph G, let $\tau(G)$ denote the smallest size of a set of vertices S which are incident to every edge of G. If G is bipartite, then*

$$\nu(G) = \tau(G).$$

We note that the inequality $\nu(G) \le \tau(G)$ trivially holds for every graph G, so the difficult part is in proving $\nu(G) \ge \tau(G)$ for bipartite graphs. For this next result, we define for a set of vertices S its neighborhood $N(S) = \{x : \exists y \in S,\ xy \in E(G)\}$.

Theorem 5.2 (Hall's Theorem)**.** *Let G be a bipartite graph with bipartition $U \cup V$. Then G has a perfect matching if and only if $|N(S)| \ge |S|$ for all $S \subseteq U$.*

The fact that the condition $|N(S)| \ge |S|$ is necessary for a perfect matching is immediate, so again the difficulty lies in proving this is sufficient for bipartite graphs.

We now move on to prove a slightly less well-known result characterizing when *arbitrary* graphs G have a perfect matching by showing that an "obvious" necessary condition is also sufficient. To this end, given a graph G we let $\mathrm{odd}(G)$ denote the number of connected components of G which have an odd number of vertices.

Theorem 5.3 (Tutte's Theorem)**.** *A graph G has a perfect matching if and only if $\mathrm{odd}(G-S) \le |S|$ for all $S \subseteq V(G)$.*

Proof. We leave the proof that $\mathrm{odd}(G-S) \le |S|$ for all S is a necessary condition as an exercise to the reader. □

Can state Tutte-Berge as analog of Konig.

TODO. Likely topics: Konig's Theorem, Hall's Theorem, stable matchings, Tutte's Theorem, Factors

5.1 Exercises

1. Prove that for every graph G we have $\nu(G) \leq \tau(G)$ where $\nu(G)$ is the size of a largest matching in G and $\tau(G)$ is the smallest size of a set of vertices S which are incident to every edge of G.

6 Connectivity and Flows

TODO. Likely topics: k-connectivity, blocks, Menger's Theorem, min cut max flow, ear decompositions

Part III

Methods

7 Probabilistic Methods

One the most important development in extremal combinatorics has been the idea of using probabilistic tools to solve extremal problems. We've already seen a few examples of this: in Theorem 3.7 we used a uniform random edge-coloring to prove exponential lower bounds on the Ramsey number $R(n)$, and in Theorem 4.17 we used a certain random graph (after deleting a few of its vertices) to prove the existence of graphs with high chromatic number and high girth, and it is hard to understate how many other breakthrough results beyond these were solved using probabilistic thinking.

Since GSU has a whole class dedicated to this topic, we reluctantly will restrain ourselves and give only a single additional application of this beautiful method in this text. We encourage anyone interested in digging further to read either the standard text by Alon and Spencer, as well as my own notes here. We now use the probabilistic method to give a general lower bound for Turán numbers of arbitrary graphs F.

Theorem 7.1. *If F is a graph with v vertices and e edges with $e \geq v$, then*

$$\mathrm{ex}(n, F) = \Omega(n^{2-\frac{v-2}{e-1}}).$$

Proof. Our proof approach uses the same "method of alterations" as in Theorem 4.17: we start with random graph $G_{n,p}$ (which we recall is the n-vertex graph obtained by keeping edge independently with probability p) and then delete "bad parts" of this random graph until it has our desired property (namely that of F-freeness).

To this end, consider $G_{n,p}$ with p a quantity to be determined later. Let X denote the number of copies of F in $G_{n,p}$. For S a set of v vertices, let 1_S be the indicator variable which is 1 if S contains a copy of F in $G_{n,p}$ and which is 0 otherwise. With this,

$$\sum 1_S \leq X \leq v! \sum 1_S,$$

since each set of v vertices contains at most $v!$ copies of F. To have $1_S = 1$, we in particular need S to contain at least e edges, so

$$\Pr[1_S = 1] \leq \sum_{k \geq e} \binom{\binom{v}{2}}{k} p^k (1-p)^{\binom{v}{2}-k} \leq v^2 2^{v^2} p^e \leq 4^{v^2} p^e.$$

In total this gives

$$\mathbb{E}[X] \leq v! \binom{n}{v} \cdot 4^{v^2} p^e \leq (4^v n)^v p^e.$$

Observe that when $p \gg n^{v/e}$, the calculation above suggests that $G_{n,p}$ will contain copies of F (at least in expectation), so $G_{n,p}$ will not work as an F-free graph for this range of p. However, we can get around this by observing that if G is a subgraph of $G_{n,p}$ obtained by deleting an edge from each copy of F in $G_{n,p}$ then G will be F-free by construction. Moreover, the number we will have $e(G) \geq e(G_{n,p}) - X$ since at most X of the original edges from $G_{n,p}$ are deleted. Using linearity of expectation gives

$$\mathbb{E}[e(G)] \geq \mathbb{E}[e(G_{n,p}) - X] \geq p\binom{n}{2} - (4^v n)^v p^e \geq \frac{1}{4} pn^2 - (4^v n)^v p^e. \tag{4}$$

At this point we want to choose p so that the above expression is maximized. Intuitively this will happen when both terms on the rightside of (4) are roughly equal to each other, i.e. when $pn^2 \approx n^v p^e$. This suggests taking $p \approx n^{\frac{2-v}{e-1}}$. And indeed, after playing around for a bit, one sees that, for example, taking $p = \frac{1}{20 \cdot 16^v} n^{\frac{2-v}{e-1}}$ and plugging it into (4) gives[10] $\mathbb{E}[e(G)] \geq \frac{1}{160 \cdot 16^v} n^{2-\frac{2-v}{e-1}}$. Because G is a (random) F-free graph with at least this many edges in expectation, there must exist some deterministic F-free graph with at least this many edges, proving the result. □

Theorem 7.1 can fail to be effective if we consider F with, say, a bunch of isolated vertices. However, a simple observation allows one to improve upon Theorem 7.1 in cases like these.

Corollary 7.2. *For every graph F with $e(F) \geq 2$ we define the 2-density*

$$m_2(F) := \max_{F' \subseteq F,\ e(F') \geq 2} \frac{e(F') - 1}{v(F') - 2}.$$

For any F with $e(F) \geq 2$ we have

$$\mathrm{ex}(n, F) = \Omega(n^{2 - \frac{1}{m_2(F)}}).$$

Proof. If $m_2(F) = 1$ then we only need to prove $\mathrm{ex}(n, F) = \Omega(n)$ which is trivial via considering an n-vertex star if $F \neq K_{1,t}$ and considering an n-vertex matching otherwise. We can thus assume $m_2(F) > 1$ from now on.

Let $F' \subseteq F$ be any subgraph obtaining the maximum in $m_2(F)$, which under the assumption of $m_2(F) > 1$ implies that $e(F') \geq v(F')$. Thus by Theorem 7.1 we have

$$\mathrm{ex}(n, F) \geq \mathrm{ex}(n, F') = \Omega(n^{2 - \frac{v(F')-2}{e(F')-1}}) = \Omega(n^{2 - \frac{1}{m_2(F)}}),$$

where here this first inequality used the fact that any F'-free graph is also F-free. □

7.1 Exercises

1. Prove that every graph G contains a bipartite subgraph G' with $e(G') \geq \frac{1}{2} e(G)$.

   (a) Using probability [2].

   (b) Without using probability [2].

   (c) Prove the stronger fact that $e(G') > \frac{1}{2} e(G)$ whenever $e(G) > 0$ [2].

2. We can use probability to give another proof of Turán's Theorem.

   (a) (Caro-Wei) Prove that if G is an n-vertex, then

$$\alpha(G) \geq \sum_{x \in V(G)} \frac{1}{\deg(x) + 1}.$$

[10] Here we use $4^{v^2} \leq 4^{ve}$ and that $e \geq 2$.

(Hint: construct a random independent set I in such a way that $\Pr[x \in I] = \frac{1}{\deg(x)+1}$) [2+].

(b) Conclude $\mathrm{ex}(n, K_r) \leq (1 - \frac{1}{r-1})\frac{n^2}{2}$.

3. Let G be a graph with m edges and N copies of a graph F.

   (a) Prove (without using probability) that G contains an F-free subgraph with at least $m - N$ edges. [1+]

   (b) Prove that if $N \geq m$ and $e(F) \geq 2$, then G contains an F-free subgraph with at least

   $$\Omega\left(\frac{m^{1+\frac{1}{e(F)-1}}}{N^{\frac{1}{e(F)-1}}}\right)$$

   edges [2].

   c What does this result apply when $G = K_n$? [1]

We note that our intended proof of (b) here uses probability to "boost" the weak deterministic bound from (a) to get a substantially stronger bound. This is a common phenomenon with the probabilistic method, as the next exercise aims to show as well.

4. Given a graph G, define the crossing number $cr(G)$ to be the minimum number of crossing pairs of edges in any embedding of G into $\mathbb{R}^2$. For example, $cr(G) = 0$ if and only if G is planar.

   (a) Prove (without using probability) that $cr(G) \geq e(G) - 3v(G)$ (Hint: use Euler's formula) [2].

   (b) (Crossing lemma) Prove that there exists some $C > 0$ such that if G is an n-vertex graph with $m \geq Cn$ edges then

   $$cr(G) = \Omega\left(\frac{m^3}{n^2}\right).$$

   [2+]

8 Regularity and Removal Lemmas

This chapter is centered around *Szemerédi's regularity lemma* (or simply *the regularity lemma* for short), which was originally a lemma proven by Szemerédi in his proof of Szemerédi's Theorem but which has since been recognized as a powerful and fundamental tool in graph theory with many applications to Turán problems, Dirac problems, and much more. We will only scratch the surface on what can be said here, and we refer the interested reader to the book "Graph Theory and Additive Combinatorics" by Zhao for a more thorough treatment.

8.1 The Regularity Lemma and its Applications

Informally, the regularity lemma says that for every graph G, there exists a partition of $V(G)$ into a bounded number of parts such that the graph between most pairs of parts "looks like" a random graph. We need some definitions to make this precise.

Definition 17. Let G be an n-vertex graph, $A, B \subseteq V(G)$ sets of (not necessarily disjoint) vertices, and $\varepsilon > 0$ a real number.

- We define $e(A, B) = |\{(x, y) \in A \times B : xy \in E(G)\}$. Note that $e(A, B)$ equals the number of edges between A and B if these sets are disjoint, and $e(A, B)$ equals twice the number of edges within A if $A = B$.

- We define the *density* of the pair (A, B) by

$$d(A, B) = \frac{e(A, B)}{|A||B|}.$$

  Note that $0 \leq d(A, B) \leq 1$ for all A, B.

- We say that the pair (A, B) is *ε-regular* if for any $X \subseteq A$ and $Y \subseteq B$ with $|X| \geq \varepsilon|A|$ and $|Y| \geq \varepsilon|B|$, we have

$$|d(X, Y) - d(A, B)| \leq \varepsilon.$$

  That is, a pair is ε-regular if all large subsets of A, B have roughly the same density as the pair itself. Note that this sort of property is what we would expect to see if we constructed a random graph on $A \cup B$ by keeping each edge independently and with probability $d(A, B)$.

- We say that a partition $V_1 \cup \cdots \cup V_m$ of $V(G)$ is an *ε-regular partition* if

$$\sum_{(V_i, V_j) \text{ which are not } \varepsilon\text{-regular}} e(V_i, V_j) \leq \varepsilon n^2.$$

.

Theorem 8.1 (Szemerédi's Regularity Lemma)**.** *For all $\varepsilon > 0$, there exists a number $M(\varepsilon)$ such that every graph G has an ε-regular partition $V_1 \cup \cdots \cup V_m$ with $m \leq M(\varepsilon)$*

We note that the number $M(\varepsilon)$ is horrendously large in terms of ε, which means that the implicit constants in any proof which uses regularity lemma will be horrendously large as well. While we will not get into it here, a result of Gowers shows that $M(\varepsilon)$ must necessarily be quite large for the theorem to hold. In a similar spirit, one can cook up constructions which show that this result is not true if we replace the condition that most edges are in ε-regular pairs with the condition that *all* edges are in ε-regular pairs.

We now give a proof of the regularity lemma. We emphasize, however, that the reader may first find it helpful to actually read through some of the applications of this result first to get a feel for these strange definitions and only then come back to review its proof.

Proof. We'll omit the proof due to limited time, though the actual argument is not actually that involved. Roughly we start with some arbitrary partition of G, iteratively if some pair V_i, V_j fails to be ε-regular then we can further partition each of V_i, V_j into a bounded number of parts which "improves" the partition in some way until eventually the resulting thing does what we want. □

We now turn to applications of the regularity lemma, which in general all go through the same three basic steps:

(1) Take an ε-regular partition $V_1 \cup \cdots \cup V_m$ for your graph G as guaranteed by the regularity lemma.

(2) "Clean" the graph G by deleting a small number of "poorly behaved" edges, e.g. by deleting all edges between any pairs (V_i, V_j) which are either not ε-regular, have low density, or which have $|V_i|$ relatively small.

(3) Solve the problem for the cleaned graph, often by invoking known results from extremal combinatorics or by making use of a "counting" lemma.

One basic and very important example of this framework comes from the following result known as the triangle removal lemma (or simply the removal lemma depending on context).

Theorem 8.2 (Triangle Removal Lemma)**.** *For all $\varepsilon > 0$, there exists a $\delta > 0$ such that if G is an n-vertex graph with at most δn^3 triangles, then G can be made triangle-free by deleting at most εn^2 edges.*

Essentially, this says that any graph which is close to being triangle-free (in the sense that it has $o(n^3)$ triangles) can be made triangle-free by deleting at most $o(n^2)$ edges.

Proof. Fix some $\varepsilon > 0$, and with plenty of foresight we define $\varepsilon' = \min\{\frac{1}{4}, \varepsilon/4\}$ and $\delta = \frac{1}{8}\frac{1}{4}(\varepsilon')^6 M(\varepsilon')^{-3}$ where $M(\varepsilon')$ is as in the statement of the regularity lemma. Let G be an n-vertex graph with at most δn^3 triangles. With our framework above in mind, we begin by applying the regularity lemma to obtain an ε'-regular partition $V_1 \cup \bigcup V_m$ and then we define our "cleaned" subgraph $G' \subseteq G$ by deleting all edges between pairs (V_i, V_j) such that either (V_i, V_j) is not ε'-regular, or $d(V_i, V_j) \leq 2\varepsilon'$, or $\min\{|V_i|, |V_j|\} \leq \varepsilon' m^{-1} n$.

Claim 8.3. *It suffices to show that the graph G' is triangle-free.*

Proof. Observe that the number of edges we deleted going from G to G' is certainly at most

$$\varepsilon' n^2 + \sum_{i,j} 2\varepsilon' |V_i||V_j| + \sum_{i:|V_i| \le \varepsilon' m^{-1} n} |V_i| n \le \varepsilon' n^2 + 2\varepsilon' + \varepsilon' n = 4\varepsilon' n^2 \le \varepsilon n^2,$$

where this first inequality used that the number of terms in the sum is at most m. As such, if we assume the hypothesis of the claim then we see that we can indeed remove at most εn^2 edges from G so that the resulting graph is triangle-free. □

Assume for contradiction that G' contains a triangle, say with each of its vertices coming from parts V_i, V_j, V_j and we emphasize that we do not require that these integers i, j, k to be distinct from each other. This implies that there is at least one edge in G' between each of these three parts, which by definition of G' implies that all of these pairs of parts are ε'-regular, have density at least $2\varepsilon'$, and that each of these parts has size at least $\varepsilon' m^{-1} n$.

Claim 8.4. *The set of vertices $V_i' \subseteq V_i$ which have at least $\varepsilon'|V_j|$ neighbors in V_j and at least $\varepsilon'|V_k|$ neighbors in V_k satisfies $|V_i'| \ge (1 - 2\varepsilon')|V_i|$.*

Proof. Indeed, if we let $X \subseteq V_i$ denote the set of vertices with less than $\varepsilon'|V_j|$ neighbors in V_j, then

$$d(X, V_j) < \frac{\varepsilon' |X||V_j|}{|X||V_j|} = \varepsilon'.$$

Because $d(V_i, V_j) \ge 2\varepsilon'$ and (V_i, V_j) is ε'-regular, this inequality is only possible if $|A| \le \varepsilon'|V_i|$. An analogous argument shows that the number of vertices of V_i with less than $\varepsilon'|V_k|$ neighbors in V_k is at most $\varepsilon'|V_i|$, proving the claim. □

Claim 8.5. *Every $x_i \in V_i'$ as defined above is contained in at least $\frac{1}{2}(\varepsilon')^3|V_j||V_k|$ triangles.*

Proof. By definition of V_i', the sets $X_j := N_{G'}(x_i) \cap V_j$ and $X_k := N_{G'}(x_i) \cap V_k$ both have at least an ε' proportion of the vertices of V_j, V_k. Using this and the fact that the pair (V_j, V_k) is ε'-regular, we find that

$$d(V_j', V_k') \ge d(V_j, V_k) - \varepsilon' \ge \varepsilon',$$

which by definition means that the number of pairs of adjacent vertices $(x_j, x_k) \in V_j' \times V_k'$ is at least $\varepsilon'|V_j'||V_k'| \ge (\varepsilon')^3|V_j||V_k|$. Since each of these pairs (x_j, x_k) forms a triangle with x_i and since each such triangle arises from at most 2 pairs (x_j, x_k), we find that the number of triangles using x_i is as stated. □

By combining these two claims, we see that for any choice of $\varepsilon' \le \frac{1}{4}$ that the number of triangles in G' is at least

$$\frac{1}{4}(\varepsilon')^3|V_i||V_j||V_k| \ge \frac{1}{4}(\varepsilon')^6 m^{-3} n^3 \ge \frac{1}{4}(\varepsilon')^6 M(\varepsilon')^{-3} n^3 = 2\delta n^3,$$

with the first inequality using that every part surviving in G' has size at least $\varepsilon' m^{-1} n$. This in turn implies that $G \supseteq G'$ contains at least $2\delta n^3$ triangles, a contradiction to our assumption that it contains at most δn^3 triangles, proving the result. □

As an aside, because our proof used the regularity lemma, the dependencies of ε and δ we obtain are very bad. There do exist regularity-free proofs of the triangle removal lemma due to Fox which gives nearly optimal bounds on these dependencies, but even these nearly optimal bounds are still quite large.

The triangle removal lemma is, in addition to simply being a nice statement, an incredibly important tool in its own right. We will see some application of this in the next subsection. For now , we continue looking at applications of the regularity lemma by proving the Erdős-Stone-Simonovits Theorem which we recall below.

Theorem 8.6. *For any graph F with at least one edge, we have*

$$\mathrm{ex}(n,F)=\left(\frac{\chi(F)-2}{\chi(F)-1}+o(1)\right)\binom{n}{2}.$$

Proof. The key observation due to Simonovits is that we can reduce our problem to studying complete r-partite graphs where every part has size t, and we let $K_{r;t}$ denote this graph.

Claim 8.7. *It suffices to prove for all $r,t\ge 2$ that*

$$\mathrm{ex}(n,K_{r;t})\le\left(\frac{r-2}{r-1}+o(1)\right)\binom{n}{2}.$$

Proof. Assume this is true and consider any graph F with $r:=\chi(F)$. The lower bound comes from considering $G=K_{r-1;\lfloor n/(r-1)\rfloor}$. For the upper bound, we observe that $F\subseteq K_{r;v(F)}$ since any $\chi(F)=r$ implies that F is r-partite, and as such it is certainly contained in a complete r-partite graph with every part of size F. As such,

$$\mathrm{ex}(n,F)\le\mathrm{ex}(n,K_{r;v(F)})\le\left(\frac{r-2}{r-1}+o(1)\right)\binom{n}{2},$$

proving the result. □

We now prove this upper bound on $\mathrm{ex}(n,K_{r;t})$. This was originally done by Erdős and Stone, though we emphasize that their proof was not based on regularity like ours is. Many of the details here are completely analogous to what we did in our proof of the triangle removal lemma, so we will a bit terser in our exposition whenever the parallels are clear.

Fix some r,t. Proving this asymptotic bound is equivalent to showing that for any $\delta>0$, we have $\mathrm{ex}(n,K_{r;t})\le(\frac{r-2}{r-1}+\delta)\binom{n}{2}$ for all sufficiently large n. To this end, fix some $\delta>0$, let $d>0$ be some small constant in terms of δ,r,t, and let $\varepsilon>0$ be a very, very small constant in terms of d,r,t (we won't specify it exactly, but we will want $\varepsilon\approx(d/2)^{rt}$).

Let G be an n-vertex graph with at least $(\frac{r-2}{r-1}+\delta)\binom{n}{2}$ edges and $V_1\cup\cdots\cup V_m$ an ε-partition of G as guaranteed by the regularity lemma. Let $G'\subseteq G$ be the subgraph defined by deleting all edges between V_i and V_j for any pair (V_i,V_j) which is either not ε-regular, or has $d(V_i,V_j)\le d$, or has $\min\{|V_i|,|V_j|\}\le\varepsilon m^{-1}n$.

As in our proof of the triangle removal lemma, we observe that the number of edges we delete when going from G to G' is at most

$$\varepsilon n^2+dn^2+\varepsilon n^2\le\frac{\delta}{2}n^2,$$

with the last step using our assumption of ε, d being sufficiently small in terms of δ. In particular, G' is an n-vertex graph with at least $(\frac{r-2}{r-1} + \delta/2)\binom{n}{2}$ edges, so for n sufficiently large Turán's Theorem guarantees that G' contains a K_r. Possibly by relabeling our parts we can assume that the r vertices of this K_r lie in parts $V_1, \ldots, V_r$ where again we allow the possibility that $V_i = V_j$ for some $i \neq j$. The existence of this K_r implies that there exist edges in G' between each of the parts V_i, V_j for all distinct $1 \leq i, j \leq r$, so by construction of G' this implies that each of these pairs (V_i, V_j) is ε-regular and has $d(V_i, V_j) \geq d$.

Intuitively at this point we will try and proceed as follows to build our copy of $K_{r;t}$: we begin by selecting some vertex $x_{1;1} \in V_1$ which has at least $(d-\varepsilon)|V_j|$ in each of the other parts, with most vertices in V_1 satisfying this property. We then pick $x_{1;2} \in V_1$ which has at least $(d-\varepsilon)|V_j \cap N(x_{1;1})|$ neighbors in each set $V_j \cap N(x_{1;1})$ which again is satisfied by most choices of vertices in V_1 provided ε is much smaller compared to d. We then pick $x_{1;3}$ to have many neighbors within the common neighborhoods $N(x_{1;1}) \cap N(x_{1;2})$ and continue in this way until we have selected vertices $x_{1;1}, \ldots, x_{1;t}$ which have many common neighbors in each of the other parts. From here we pick some $x_{2;1}$ in the intersection of this common neighborhood and V_2 and proceed in a similar way until we have eventually constructed a full copy of $K_{r;t}$. The following gives a formal framework for this approach to work.

Claim 8.8. *For n sufficiently large, we have that if $\widetilde{V}_1, \subseteq, \widetilde{V}_r$ are subsets of $V_1, \ldots, V_r$ such that $(d-\varepsilon)^t|\widetilde{V}_i| \geq 2r\varepsilon|V_i|$ for all i, then for all i and integers $1 \leq s \leq t$ there are at least $2^{-s}|\widetilde{V}_i|^s$ tuples $(x_1, \ldots, x_s) \in \widetilde{V}_i^s$ of distinct vertices such that $|\bigcap_{s'=1}^{s} N(x_{s'}) \cap \widetilde{V}_j| \geq (d-\varepsilon)|\widetilde{V}_j|$ for all $j \neq i$*

Proof. We prove the result by induction on s, the base case $s = 0$ being trivial. Let $(x_1, \ldots, x_{s-1})$ be an arbitrary tuple satisfying the conditions for $s-1$, let $\widetilde{V}_j^* = \bigcap_{s'=1}^{s-1} N(x_{s'}) \cap \widetilde{V}_j$ and let X_j denote the set of $x \in \widetilde{V}_i$ such that $|N(x) \cap \widetilde{V}_j^*| < (d-\varepsilon)|\widetilde{V}_j|$. Then

$$d(X_j, \widetilde{V}_j^*) < d - \varepsilon \leq d(V_i, V_j) - \varepsilon.$$

Because (V_i, V_j) is an ε-regular pair and $|\widetilde{V}_j^*| \geq (d-\varepsilon)^{s-1}|\widetilde{V}_j|\varepsilon|V_j| \geq \varepsilon|V_j|$ by hypothesis, we must have $|X_j| \geq \varepsilon|V_i|$. Using this and the requirement that our tuples have distinct vertices, we find that the number of choices for x which we can append to $(x_1, \ldots, x_{s-1})$ while satisfying the conditions of the claim is at least

$$|\widetilde{V}_i \setminus (\bigcup X_j)| - s + 1 \geq |\widetilde{V}_i| - (r-1)\varepsilon|V_i| - t \geq |\widetilde{V}_i| - r\varepsilon|V_i| \geq \frac{1}{2}|\widetilde{V}_i|,$$

with the middle inequality holding for n sufficiently large in terms of ε since $|V_i| \geq \varepsilon m^{-1} n \geq \varepsilon M(\varepsilon)^{-1} n$. Since we inductively assumed the number of choices for $(x_1, \ldots, x_{s-1})$ was at least $2^{1-s}|\widetilde{V}_i|^{s-1}$ we conclude the result. $\square$

To prove the result, we first apply this claim with $\widetilde{V}_j = V_j$ for all j to find a tuple $(x_{1;1}, \ldots, x_{1;t}) \in V_1^t$ satisfying these conditions. Iteratively given that we have constructed $(x_{i';1}, \ldots, x_{i';t})$ for all $1 \leq i' < i \leq r$, we apply the claim with $\widetilde{V}_j = V_j \bigcap_{1 \leq i' < i,\ 1 \leq s \leq t} N(x_{i';s})$ for $j \geq i$ and $\widetilde{V}_j = V_j$ for $j < i$ to find a tuple $(x_{i;1}, \ldots, x_{i;t}) \in \widetilde{V}_i$. Note that iteratively we always apply the claim after assuming $|\widetilde{V}_j| \geq (d-\varepsilon)^{rt}|V_j|$, so the hypothesis of the claim will hold provided ε is sufficiently small in terms of d, r, t. We conclude that this process terminates with distinct vertices such that $x_{i;s} \sim x_{i';s'}$ whenever $i < i'$ by construction, giving our copy of $K_{r;t}$ as desired. $\square$

As an aside, we note that our proof in fact shows the somewhat stronger fact that having $e(G) \geq (\frac{r-2}{r-1} + \delta)\binom{n}{2}$ implies that G contains $\Omega_\delta(n^{rt})$ copies of $K_{r;t}$).

The exact mechanics of our proofs for both the triangle removal lemma and the Erdős-Stone-Simonovits Theorem are very similar to each other, in that they both rely on showing that if there exists a set of r parts whose pairs are all ε-regular and have high density, then one can find many copies of K_r. Arguments of this form are very common with the regularity lemma, and as such it can be useful to record facts like these into so-called "counting lemmas", a general version of which is the following.

Lemma 8.9 (Graph Counting Lemma). *For every graph F with vertex set $\{v_1, \ldots, v_r\}$ and for every real number $\delta > 0$, there exists some $\varepsilon > 0$ such that the following holds: if G is a graph and $V_1, V_2, \ldots, V_{v(F)} \subseteq V(G)$ are such that (V_i, V_j) is ε-regular with $d(V_i, V_j) \geq \delta$ whenever $v_iv_j \in E(F)$, then the number of homomorphisms ϕ from F to G with $\phi(v_i) \in V_i$ is at least*

$$(1-\delta) \prod_{v_iv_j \in E(F)} (d(V_i, V_j) - \delta) \prod_i |V_i|.$$

To be clear, this lemma only guarantees many homomorphisms of F and not necessarily many copies of F. However, the number of homomorphisms of F which are not injective is at most $O(v(G)^{v(F)-1})$, so this result guarantees many injective homomorphisms (and hence copies of F) whenever $|V_i| = \Omega(v(G))$ for all i. The proof of the graph counting lemma is spiritually similar to the proofs we have done up to this point, and as such we leave its proof as an exercise to the reader.

8.2 Applications of the Removal Lemma

We now discuss applications of the triangle removal lemma, which we recall says that any graph with $o(n^3)$ triangles can be made triangle-free by deleting at most $o(n^2)$ edges. We begin with a Turán type problem. To this end, given graphs H and F we define the *generalized Turán number* $\mathrm{ex}(n, H, F)$ to be the maximum number of copies of H in an n-vertex F-free graph. Note that $\mathrm{ex}(n, K_2, F) = \mathrm{ex}(n, F)$.

Theorem 8.10. *We have $\mathrm{ex}(n, K_3, K_4 - e) = o(n^2)$ where $K_4 - e$ denotes the graph obtained from K_4 after deleting an edge.*

That is, every n-vertex graph where every edge is contained in at most one triangle has $o(n^3)$ triangles. This is equivalent to saying that every n-vertex graph where every edge is contained in exactly one triangle has $o(n^3)$ triangles, which is the most common way this theorem appears in the literature.

Proof. Let G be any n-vertex $(K_4 - e)$-free graph. By definition every edge of G is contained in at most one triangle, implying that the number of triangles is at most $e(G)/3 = O(n^2) = o(n^3)$.

By the triangle-removal lemma one can delete $o(n^2)$ edges of G to make it triangle-free. But by definition of G, each edge removed destroys at most one triangle in G, implying that G must have had $o(n^2)$ triangles to begin with, proving the result. $\square$

This bound of $\mathrm{ex}(n, K_3, K_4 - e) = o(n^2)$ is best possible in that there exists a construct showing that $\mathrm{ex}(n, K_3, K_4 - e) \geq n^{2-o(1)}$ Maybe see exercises for more on this.

We now give some applications of the removal lemma to areas outside of combinatorics. We begin with the original motivation for the regularity lemma, namely in determining how large a subset $A \subseteq [n]$ can be if it contains no k-term arithmetic progression, i.e. no k distinct integers $a_1, \ldots, a_k \in A$ such that $a_{i+1} - a_i = a_{j+1} - a_j$ for all i, j. The simplest non-trivial case of $k = 3$ was originally solved by Roth using Fourier analysis. A substantially simpler proof can be given using the removal lemma.

Theorem 8.11 (Roth's Theorem). *If $A \subseteq [n]$ contains no 3-term arithmetic progression, then $|A| = o(n)$.*

Again the bound of $|A| = o(n)$ is best possible here as for all k there exist sets $A \subseteq [n]$ without k-AP's with size $|A| \geq n^{1-o(1)}$.

Proof. Let $A \subseteq [n]$ be such that it contains no 3-term arithmetic progression, which we crucially observe is equivalent to saying that no distinct $x, y, z \in A$ satisfy $x + y = 2z$ since in this case x, z, y would be a progression with common difference $z - x = y - z$. We now wish to construct a graph G_A whose edges are defined based on A such that G_A inherits some nice properties because A is 3-AP free. After a lot of thought, one might be led to the following idea for constructing G_A:

- The graph $G_A$is tripartite with parts $V_1 = [n],\ V_2 = [2n]$, and $V_3 = [3n]$,
- We have $v_1 \in V_1$ adjacent to $v_2 \in V_2$ if and only if there exists $a \in A$ with $v_1 + a = v_2$,
- We have $v_2 \in V_2$ adjacent to $v_3 \in V_3$ if and only if there exists $a \in A$ with $v_2 + a = v_3$, and
- We have $v_1 \in V_1$ adjacent to $v_3 \in V_3$ if and only if there exists $a \in A$ with $v_1 + 2a = v_3$.

We emphasize that the adjacency condition for V_1, V_3 is defined differently compared to the other cases. Crucially, this graph does inherit nice properties whenever A is 3-AP free.

Claim 8.12. *If $A \subseteq [n]$ is 3-AP free, then $(v_1, v_2, v_3) \in V_1 \times V_2 \times V_3$ is the vertex set of a triangle in G_A if and only if $v_2 = v_1 + a$ and $v_3 = v_1 + 2a$ for some $a \in A$.*

Proof. It is straightforward to check that every triple of vertices $(v_1, v_1 + a, v_1 + 2a)$ with $a \in A$ is a triangle. Assume now that (v_1, v_2, v_3) forms a triangle in G_A. By definition this means $v_2 - v_1 := a \in A,\ v_3 - v_2 := b \in A$, and $v_3 - v_1 := 2c$ for some $c \in A$. As such we have

$$a + b = (v_2 - v_1) + (v_3 - v_2) = v_3 - v_1 = 2c.$$

Because A is 3-AP free, this equality is only possible if at least two of a, b, c are equal to each other, but one can check that this is only possible if $a = b = c$, proving the claim. $\square$

From this claim we conclude that the number of triangles in G_A is exactly $|A|n$ when A is 3-AP free since a triangle is uniquely identified by picking $v_1 \in V_1$ and the $a \in A$ such that $v_2 = v_1 + a$. This claim also implies that every edge of G_A is contained in at most one triangle since, for example, any edge v_1v_2 with $v_1 \in V_1, v_2 \in V_2$ can only be in a triangle with $v_3 = v_1 + 2(v_2 - v_1) \in V_3$. By Theorem 8.10 we conclude that

$$|A|n = o(n^2),$$

and hence that $|A| = o(n)$, proving the result. $\square$

It is perhaps tempting giving the simplicity of this argument to try and use some sort of removal lemma to try and prove Szemerédi's Theorem that sets $A \subseteq [n]$ without k-AP's have $|A| = o(n)$, but this turns out to be substantially harder to do with perhaps the "simplest" such proof being those that rely on the difficult machinery of hypergraph removal lemmas.

The last application we consider is from property testing. In this setup, we want to quickly determine whether a given graph G either has some desired property or if it is far from having this property. To this end, we say that an n-vertex graph G is *ε-close* to having a property $\mathcal{P}$ if there exist sets of edges $E, E' \subseteq K_n$ with $|E|, |E'| \le \varepsilon n^2$ such that $G + E - E'$ has property $\mathcal{P}$ and we say that G is *ε-far* otherwise. That is, G is ε-far from $\mathcal{P}$ if we can not get G to satisfy $\mathcal{P}$ even after changing up to $2\varepsilon n^2$ of its edges. Determining precisely whether a graph satisfies a property or is far from it can take quite a bit of time if n is large. Remarkably, the property of triangle-freeness can be tested with arbitrarily high probability after checking only $O(1)$ vertices of G.

Theorem 8.13. *For all $\varepsilon, c > 0$ there exists a randomized algorithm which runs in time $O_{\varepsilon,c}(1)$ which correctly determines whether a given graph G is either triangle-free or ε-far from being triangle-free with probability at least $1 - c$.*

Proof. Let C be some large (but fixed) integer depending on ε, c to be determined later. Our algorithm goes as follows: we uniformly at random pick three vertices $v_1, v_2, v_3 \in V(G)$ and test if these vertices form a triangle. We repeat this process independently for a total of C times. If none of these vertices form a triangle then we output that G is triangle-free, and otherwise if some v_1, v_2, v_3 form a triangle then we output that G is ε-far from being triangle-free.

This algorithm always correctly identifies that G is triangle-free, so it remains to check that it correctly identifies G as being ε-far from triangle-free with high probability. And indeed, observe that G being ε-far from triangle-free means that G can not be made triangle-free by removing at most εn^2 edges. The contrapositive of the triangle removal lemma then implies that G must contain at least δn^3 triangles for some δ depending only on ε. As such, the probability that three uniform random vertices of G form a triangle is at least δ, and hence the probability that the algorithm above fails to find any triangle in its C trials is at most

$$(1 - \delta)^C,$$

and this quantity can be made less than c by taking C sufficiently large in terms of ε, c, proving the result. $\square$

8.3 Variants

There are many variants of both the regularity lemma and the removal lemma. For example, there exist versions of the regularity lemma which guarantee that each part in the partition has nearly the same size, as well as variants of the regularity lemma for settings other than graphs such as arithmetic regularity lemmas. Similarly the removal lemma can be extended to hold for arbitrary graphs (as we discuss in the exercises), as well as in settings where we care about induced copies. We will not discuss these further here and instead refer the reader again to the great book by Yufei Zhao on this topic.

8.4 Exercises

1. Prove that if G is an n-vertex graph with less than $\varepsilon^3 n^2$ edges, then $V_1 = V(G)$ is an ε-regular partition with only one part. Because of this, ε-regular partitions are only interesting and useful in the case when G has $\Theta(n^2)$ edges [1].

2. Prove the graph counting lemma [2-].

3. Using the graph counting lemma, prove the graph removal lemma: for every graph F and $\varepsilon > 0$ there exists some $\delta > 0$ such that if G is an n-vertex graph with at most $\delta n^{v(F)}$ copies of F, then G can be made F-free by removing at most εn^2 edges of G [2-].

Probably add more, eg walking through Behrand or Ramsey-Turán.

9 Linear Algebra Methods

Roughly speaking, the *linear algebra method* in combinatorics works as follows:

1 Associate a "linear algebraic object" M to your problem (e.g. a matrix or a list of vectors).

2. Determine algebraic information about M (e.g. its rank, eigenvalues, eigenvectors),

3. Use this algebraic information to conclude something about your original problem.

The linear algebra method applies to a broad range of problems. We only scratch the surface here, and we refer the reader to books by Babai and Frankl and by Matoušek for a more thorough treatment of this versatile method.

9.1 Introduction to Spectral Graph Theory

Within the context of graph theory, perhaps the most natural linear algebraic object to consider is the adjacency matrix. To this end, given a graph G we define its adjacency matrix $A(G)$ to be the symmetric matrix whose rows and columns are indexed by $V(G)$ where $A(G)_{i,j} = 1$ if $i \sim j$ in G and $A(G)_{i,j} = 0$ otherwise. We write A instead of $A(G)$ whenever G is clear from context.

A priori, A is simply a convenient way to encode the graph G, and as such there is no reason to really study A as a linear operator. Surprisingly, the algebraic properties of A contains a tremendous amount of combinatorial information about G. Because of this, there is a large area known as *spectral graph theory* which centers around studying algebraic properties of both A as well as many other types of matrices that can be associated to graphs. We will only get a glimpse of this area here and refer the reader to later chapter for more on this.

We begin with a classical connection between A and combinatorial properties of G, namely that of closed walks. For this, we recall that a sequence of vertices $(w_1, \ldots, w_{k+1})$ of a graph G is called a *walk* of length k if $w_{i+1} \in N(w_i)$ for all $1 \le i \le k$, and we say that this walk is *closed* if $w_1 = w_{k+1}$. In what follows, we use the standard linear algebra fact that every real symmetric matrix (such as A) with n rows and columns has n real eigenvalues as well as a orthonormal basis of eigenvectors.

Lemma 9.1. *If G is an n-vertex graph and if $\lambda_1, \ldots, \lambda_n$ are the eigenvalues of its adjacency matrix A, then the number of closed walks of G of length k equals $\sum_i \lambda_i^k$.*

Proof. We first observe that if G is a graph and $u, v \in V(G)$, then the number of walks of length k from u to v is $A^k_{u,v}$. Indeed, by definition of matrix multiplication, we have

$$A^k_{u,v} = \sum A_{uw_1} \cdots A_{w_{k-1}v},$$

where the sum ranges over all sequences of vertices $w_1, \ldots, w_{k-1}$. Each term of this sum will be 1 if this sequence defines a walk and will be 0 otherwise, showing that $A^k_{u,v}$ is the desired amount.

From this observation, we see that the number of closed walks of length k is exactly

$$\sum_{u \in V(G)} A_{u,u}^k = \mathrm{Tr}(A^k) = \sum_i \lambda_i^k,$$

where here the first equality used the definition of the trace of a matrix, and the second equality used both that the trace of a square matrix equals the sum of its eigenvalues and that raising a square matrix to a power k raises all of its eigenvalues to the power k as well. ☐

Corollary 9.2. *If G is an n-vertex graph and if $\lambda_1, \ldots, \lambda_n$ are the eigenvalues of its adjacency matrix A, then*

$$2e(G) = \sum_i \lambda_i^2.$$

Proof. Observe that (w_1, w_2, w_3) is a closed walk of G of legnth 2 if and only if it is of the form (u, v, u) with $uv \in E(G)$. It follows that the number of closed walks of length 2 is exactly $2e(G)$ (since there is one for each orientation of each edge), giving the result by the previous lemma. ☐

Results allows us to use algebraic information about A to determine combinatorial information about G and vice versa as the next result shows.

Corollary 9.3. *If G is a non-empty graph and if $\lambda_{\max}, \lambda_{\min}$ are the largest and smallest eigenvalues of A, then $\lambda_{\max} > 0 > \lambda_{\min}$.*

Proof. By Lemma 9.1 (or simply by definition of A and the trace), we have that

$$\sum \lambda_i = 0.$$

On the other hand, we have

$$\sum \lambda_i^2 = 2e(G) > 0.$$

These two statements are only possible if there exists some eigenvalue which is positive and some eigenvalue which is negative, proving the result. ☐

The statements above hold for arbitrary graphs, but more can be said for particular kinds of graphs. In particular, spectral graph theory tends to be at its strongest for regular graphs in part due to the following key observation.

Lemma 9.4. *If G is a d-regular graph, then the all 1's vector $\mathbf{1}$ is an eigenvector of A with eigenvalue equal to d.*

Proof. For all $u \in V(G)$ we have

$$(A\mathbf{1})_u = \sum_v A_{u,v}\mathbf{1}_v = \sum_{v \in N(u)} 1 = d = d\mathbf{1}_u,$$

proving that $A\mathbf{1} = d\mathbf{1}$. ☐

One famous application of spectral graph theory to regular graphs comes from Hoffman's bound for the independence number of G.

Theorem 9.5 (Hoffman's Ratio Bound)**.** *Let G be a non-empty n vertex d-regular graph and A its adjacency matrix. Then*

$$\frac{\alpha(G)}{n} \leq \frac{-\lambda_{\min}}{d - \lambda_{\min}},$$

where $\lambda_{\min}$ denotes the smallest eigenvalue of A.

Proof. Let I be an independent set of size $\alpha = \alpha(G)$ and let x be the vector indexed by $V(G)$ with $x_i = 1$ if $i \in I$ and $x_i = 0$ otherwise. Observe that because I is an independent set, we have

$$x^T Ax = \sum_{i,j} x_i A_{i,j} x_j = \sum_{i,j \in I} A_{i,j} = 0.$$

Let $y_1, \ldots, y_n$ be an orthonormal eigenbasis for A with eigenvalues $\lambda_1, \ldots, \lambda_n$. Since G is regular, the all 1's vector $\mathbf{1}$ is an eigenvector with eigenvalue equal to d, so we can assume $y_1 = \mathbf{1}/\sqrt{n}$ and $\lambda_1 = d$. Writing $x = \sum c_i y_i$ for some real numbers c_i, we see that

$$\alpha = x^T x = \sum c_i^2,$$

and

$$\alpha/\sqrt{n} = \langle x, v_1 \rangle = c_1.$$

Putting all of this together, we find

$$\begin{aligned} 0 = x^T Ax = x^T \sum c_i \lambda_i v_i = \sum c_i^2 \lambda_i &= (\alpha^2/n)d + \sum_{i \neq 1} c_i^2 \lambda_i \\ &\geq (\alpha^2/n)d + \sum_{i \neq 1} c_i^2 \lambda_{\min} = (\alpha^2/n)d + (\alpha - \alpha^2/n)\lambda_{\min}. \end{aligned}$$

Dividing both sides by α and rearranging gives

$$\alpha(\lambda_{\min} - d)/n \geq \lambda_{\min}.$$

Dividing both sides by $\lambda_{\min} - d$ (which is negative because $\lambda_{\min} < 0$ since G is non-empty) gives the result. $\square$

Hoffman's Ratio Bound is quite effective for a number of graphs. In particular, one can use this to prove the Erdős-Ko-Rado Theorem, which is the fundamental theorem of extremal set theory.

Theorem 9.6 (Erdős-Ko-Rado)**.** *Let $\mathcal{F} \subseteq \binom{[n]}{r}$ be a collection of r-element subsets of $[n]$ which is intersecting, i.e. which is such that $F \cap F' \neq \emptyset$ for all $F, F' \in \mathcal{F}$. If $n \geq 2r$, then*

$$|\mathcal{F}| \leq \binom{n-1}{r-1}.$$

Note that this result is best possible by considering $\mathcal{F}$ to consist of all sets containing the element 1.

Sketch of Proof. Define an auxiliary graph G which has vertex set $\binom{[n]}{r}$ where we have $F \sim F'$ if and only if $F \cap F' = \emptyset$. From this, we see that a family $\mathcal{F}$ is intersecting if and only if it is an independent set of G. One can show that G has $\binom{n}{r}$ vertices, that it is regular with degree $\binom{n-r}{r}$, and (less trivially via using $n \geq 2r$) that the smallest eigenvalue of its adjacency matrix equals $-\frac{r}{n-r}\binom{n-r}{r}$. In total this implies that any intersecting family $\mathcal{F}$ satisfies

$$|\mathcal{F}| \leq \alpha(G) \leq \binom{n}{r} \frac{\frac{r}{n-r}\binom{n-r}{r}}{\binom{n-r}{r} - \frac{r}{n-r}\binom{n-r}{r}} = \frac{r}{n-r}\binom{n}{r} = \binom{n-1}{r-1},$$

proving the result. $\square$

9.2 Beyond the Adjacency Matrix

While the adjacency matrix is perhaps the most natural matrix to associate to a graph G, there are in fact many different types of matrices that could be considered, each of which have their own sets of advantages and disadvantages. In particular, many results and proofs which work for the adjacency matrix continue to work word for word for a slightly broader class of matrices which can sometimes be advantages to consider. To illustrate this fact, we consider another classical result relating the eigenvalues of A to combinatorial properties of G.

Lemma 9.7. *For any graph G, the largest eigenvalue $\lambda_{\max}$ of the adjacency matrix A satisfies*

$$\lambda_{\max} \leq \Delta(G).$$

Proof. Let x be an eigenvector of A corresponding to $\lambda_{\max}$ and let $v \in V(G)$ be such that $|x_v|$ is maximized. Then by our definitions, we have

$$|\lambda_{\max} x_v| = |(Ax)_v| = |\sum_u A_{v,u} x_u| \leq \sum_{u \sim v} |x_u| \leq \deg(v)|x_v| \leq \Delta(G)|x_v|.$$

This shows $|\lambda_{\max}| \leq \Delta$, proving the result. $\square$

Examining this proof, we see that we hardly used any of the properties of A in our argument. In particular, word for word the same argument gives the following.

Lemma 9.8. *Let G be a graph and M any symmetric matrix such that $M_{u,v} = \pm 1$ if $uv \in E(G)$ and $M_{u,v} = 0$ otherwise. Then the largest eigenvalue $\lambda_{\max}$ of M satisfies*

$$\lambda_{\max} \leq \Delta(G).$$

A priori it's not clear whether Lemma 9.8 is actually interesting or if it is just a generalization for generalization's sake. Surprisingly, this result plays a key role in a beautiful proof of Huang's solving a 30 year problem known as the sensitivity conjecture.

Theorem 9.9 (Huang [?])**.** *Let Q_n be the hypercube graph on 2^n vertices. If $V \subseteq V(Q_n)$ is a subset of size $2^{n-1}+1$, then the induced subgraph $Q_n[V]$ has maximum degree at least $\sqrt{n}$.*

This result is sharp in several ways. First, it is easy to find subsets of size 2^{n-1} such that $Q_n[V]$ is the empty graph, so in order to get any non-trivial lower bound on the maximum degree one needs V to have size at least $2^{n-1}+1$. Second, Chung et. al. [?] proved that there exist choices of V such that $Q_n[V]$ has maximum degree $\lceil \sqrt{n} \rceil$, so this bound is essentially best possible.

It was shown by Gotsman and Linial [?] that proving a result of this form is equivalent to showing that two notions of "sensitivity" for Boolean functions are equivalent, which led to a great deal of interest in resolving it. Nevertheless, it remained unanswered for 30 years until Huang came up with the following remarkable proof.

The key idea is to define the $2^n \times 2^n$ matrix B_n recursively by

$$B_0 = [0], \qquad B_n = \begin{bmatrix} B_{n-1} & I \\ I & -B_{n-1} \end{bmatrix},$$

where here I denotes the identity matrix of dimension 2^{n-1}. Observe that if the negative sign in the definition of B_n wasn't there, then this would just define the adjacency matrix of Q_n. Thus this is a sort of "twisted adjacency matrix" which has -1's in some of the positions where there are usually 1's. This choice of signings turns out to spread out the spectrum of B_n in a nice way.

Lemma 9.10. *The spectrum of B_n consists of $\pm\sqrt{n}$ each occurring with multiplicity 2^{n-1}.*

Proof. It is straightforward to prove by induction that $B_n^2 = nI$, which implies that every eigenvalue λ of B_n satisfies $\lambda^2 = n$. Thus $\sigma(B_n)$ consists of $\pm\sqrt{n}$, and each must appear with equal multiplicity because $\mathrm{Tr}(B_n) = 0$. □

We will also need a basic fact from linear algebra.

Theorem 9.11 (Cauchy interlacing theorem)**.** *Let B be a real symmetric $n \times n$ matrix and C an $m \times m$ principal sumbmatrix of B with $m \le n$. If B has eigenvalues $\lambda_1 \ge \cdots \ge \lambda_n$ and C has eigenvalues $\mu_1 \ge \cdots \ge \mu_m$, then for all i*

$$\lambda_i \ge \mu_i \ge \lambda_{i+n-m}.$$

Shockingly, we have everything we need for our proof of Huang's Theorem.

Proof of Theorem 9.9. Let $B = B_n$ be as described above. Let $V \subseteq V(Q_n)$ be any subset of size $2^{n-1}+1$ and let C be the submatrix of B indexed by the rows and columns corresponding to B. Let $G = Q_n[V]$. Observe that C satisfies the conditions for M of Lemma 9.8 since B is a (symmetrically) signed version of the adjacency matrix. By Lemma 9.8, the Cauchy interlacing theorem, and the previous lemma, we conclude that

$$\Delta(G) \ge \lambda_1(C) \ge \lambda_{2^{n-1}}(B) = \sqrt{n},$$

proving the result. □

9.3 Beyond Matrices

The linear algebra method extends far beyond just using eigenvalues of matrices to solve problems. To illustrate this, we briefly look at what is perhaps the most famous application of the linear algebra method, though admittedly this will require us to leave the world of graph theory and enter the related world of extremal set theory.

Consider the following (somewhat whimsical) setup. The city of Oddtown has a number of clubs, each of which follows the following odd set of rules: each club must have an odd number of people, and every two distinct clubs must have an even number of people in common.

The main question now becomes: if Oddtown has n people, what's the maximum number of clubs it can have? Equivalently, if $\mathcal{F} = \{F_1, F_2, \ldots, F_m\} \subseteq 2^{[n]}$ is a set system such that $|F_i|$ is odd for all i and such that $|F_i \cap F_j|$ is even for all $i \neq j$, then what is the maximum size of $\mathcal{F}$?

A very simple construction is to take $F_i = \{i\}$ for all i, which trivially satisfies the stated conditions. However, it's far from the only construction. For example, if n is even one can also take each F_i to be either $\{i\}$ or $[n] \setminus \{i\}$, and there are many, many more constructions achieving a bound of n (in fact, there's close to 2^{n^2} non-isomorphic constructions due to Szegedy [**?**, Exercise 1.1.14]).

Given all of these constructions, it seems plausible that (1) the true answer is indeed n, and (2) proving this might be difficult (since we have to come up with an argument that somehow deals with all of these constructions in a unified way). Fortunately, the linear algebra method manages to give a unified approach for all of these constructions in an extremely elegant way. More generally, if a given problem has many distinct looking extremal constructions, then it is often the case that the linear algebra method can come in handy.

Theorem 9.12 (Oddtown)**.** *Let $\mathcal{F} \subseteq 2^{[n]}$ be a set system such that $|F|$ is odd for all $F \in \mathcal{F}$ and such that $|F \cap F'|$ is even for all $F \neq F' \in \mathcal{F}$. Then $|\mathcal{F}| \leq n$.*

Proof. Given a set $F \subseteq [n]$, define its characteristic vector $\chi_F \in \mathbb{F}_2^n$ by having $(\chi_F)_i = 1$ if $i \in F$ and $(\chi_F)_i = 0$ otherwise. Note crucially that for any F, F', the dot product satisfies

$$\langle \chi_F, \chi_{F'} \rangle = |F \cap F'| \mod 2.$$

We claim that $\{\chi_F : F \in \mathcal{F}\}$ is a set of linearly independent vectors. Indeed, say we had

$$\sum_{F \in \mathcal{F}} \lambda_F \chi_F = 0.$$

Take any $F' \in \mathcal{F}$ and apply the dot product on both sides to get

$$\sum_{F \in \mathcal{F}} \lambda_F \langle \chi_F, \chi_{F'} \rangle = 0.$$

By the observation above and the hypothesis of the theorem, we see $\langle \chi_F, \chi_{F'} \rangle = 0$ if $F \neq F'$ and that $\langle \chi_{F'}, \chi_{F'} \rangle = 1$. Thus the above says $\lambda_{F'} = 0$, and as $F' \in \mathcal{F}$ was arbitrary, we conclude that these vectors are indeed linearly independent.

Since we have $|\mathcal{F}|$ linearly independent vectors in $\mathbb{F}_2^n$, we must have $|\mathcal{F}| \leq n$, giving the result. □

While the above technically is a proof without the use of matrices, we note that one can write an essentially equivalent proof in the language of matrices. However, for many generalizations of oddtown, the most natural way to generalize this argument is through the language of vectors (with these vectors typically being some set of low degree polynomials). We will explore this further in the exercises.

9.4 Exercises

Throughout this and maybe earlier we define the spectrum $\sigma(M)$ of a real symmetric matrix M to be the multiset of eigenvalues of A and we let $\lambda_{\max}, \lambda_{\min}$ denote the largest and smallest eigenvalues of A.

1. Prove that if G is connected and has diameter d, then A has at least $d+1$ eigenvalues [2].

2. Prove that if G is a graph with average degree $\bar{d}$, then $\lambda_{\max} \geq \bar{d}$ [2-].

3. (Wilf's Theorem) Prove that if G is a graph, then $\chi(G) \leq \lambda_{\max}+1$; note that by Lemma 9.7 this bound is always at least as strong as the classic bound $\chi(G) \leq \Delta(G)+1$ (Hint: prove this by induction on $v(G)$ via using the previous problem) [2+]. I can't really ask this without recalling the Raleigh quotient.

4. (Hoffman's Bound for the Chromatic Number) Prove that if G is a graph, then

$$\chi(G) \geq 1 - \frac{\lambda_{\max}}{\lambda_{\min}}.$$

   Note that this result is immediate from Hoffman's Ratio Bound if G is d-regular (assuming the easy to prove fact that $d = \lambda_{\max}$), so the difficulty is in proving this for non-regular graphs [2+].

5. We say that a matrix M has spectrum symmetric about 0 if the number of eigenvalues it has equal to λ is the same as the number of eigenvalues it has equal to $-\lambda$ for all λ.

   Prove that a graph G is bipartite if and only if the spectrum of its adjacency matrix is symmetric about 0 [2].

6. Prove that if A is the adjacency matrix of $K_{s,t}$, then $\sigma(A)$ has eigenvalues equal to $\sqrt{st}$ and $-\sqrt{st}$ with the rest equal to 0 [1+].

7. Prove that there exist two graphs G_1, G_2 with adjacency matrices A_1, A_2 such that $\sigma(A_1) = \sigma(A_2)$ and such that G_1 is connected while G_2 is not connected (Hint: there exist examples where G_1, G_2 have 5 vertices each) [2-].

   In general, two graphs with $\sigma(A_1) = \sigma(A_2)$ are called *cospectral*. Such graphs are important in spectral graph theory since they tell us the limitations of what can be determined by the spectrum of the adjacency matrix. For example, this result shows that one can not determine whether G is connected or not from $\sigma(A)$ alone.

* * *

8. There are various ways to generalize Hoffman's bound, here's one direction which changes how we measure the "size" of an independent set. Given a graph G, a vector x indexed by $V(G)$, and a set of vertices I, define $|I|_x = \sum_{i\in I} x_i^2$, and define $\alpha_x(G) = \max_I |I|_x$ where I ranges over all independent sets of G.

   Prove that if G is a graph and if M is a (not necessarily symmetric) matrix with rows and columns indexed by $V(G)$ such that $M_{u,v} = 0$ whenever $u \not\sim v$ and such that M has a basis of eigenvectors. If $\lambda_{\min}$ is the smallest eigenvalue of M, and if x is a unit eigenvector of M with eigenvalue $\lambda > \lambda_{\min}$, then

   $$\alpha_x(G) \leq \frac{-\lambda_{\min}}{\lambda_{\min} - \lambda}$$

   [1+].

   We note that this result can be used to prove a variant of the Erdős-Ko-Rado theorem, see reference.

9. Determine a larger class of matrices for which the results from the first couple of problems hold, eg diameter.

Part IV

Bonus Topics

10 Hypergraphs

TODO. Likely topics: generalized KST, codegree arguments and loose cycle Turán problems, Turán densities exist and supersaturation, Fisher's inequality, hypergraph ramsey, Kruskal-Katona

11 Random Graphs

TODO. Likely topics: thresholds, connectivity, spreadness theorems

12 Planar Graphs

TODO. Likely topics: Euler's formula, Wagner's Theorem characterizing planar graphs, 5-color theorem, minors.

13 Spectral Graph Theory

TODO. Likely topics: adjacency matrix, Laplacian matrix, matrix-tree theorem, Cheeger inequality, expanders

14 Advanced Methods

TODO. Likely topics: entropy, hypergraph containers, spreadness, absorption, homomorphism counting

Note: many of these topics would be covered in exactly the same way as in my notes here.